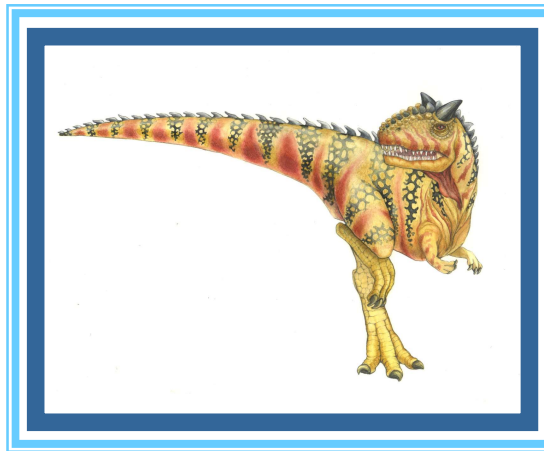


Chapter 3: Processes





Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication





Objectives

- To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication
- To explore interprocess communication using shared memory and message passing
- To describe communication in client-server systems





Process Management

- ❑ A process can be thought of as a program in execution.
- ❑ A process will **need certain resources**—such as CPU time, memory, files, and I/O devices—to accomplish its task. These resources are allocated to the process either when it is created or while it is executing.
- ❑ A process is the unit of work in most systems.
- ❑ Two types of Processes
 - ❑ **operating-system processes** execute system code, and
 - ❑ **user processes** execute user code.
- ❑ All these processes may **execute concurrently**.
- ❑ The operating system is **responsible for**
 - ❑ the **creation and deletion** of both user and system processes;
 - ❑ The **scheduling** of processes; and
 - ❑ the **provision of mechanisms** for synchronization,
 - ❑ **communication**, and
 - ❑ **deadlock handling for processes**.





Process Concept

- An operating system executes a variety of programs:
 - Batch system – **jobs**
 - Time-shared systems – **user programs** or **tasks**
- Textbook uses the terms **job** and **process** almost interchangeably
- **Process** – a program in execution; process execution must progress in sequential fashion
- A program by itself is not a process.
- **A program is a passive entity**, such as a file containing a list of instructions stored on disk (often called an executable file).
- In contrast, **a process is an active entity**, with a program counter specifying the next instruction to execute and a set of associated resources.
- **A program becomes a process** when an executable file is loaded into memory.






- Multiple parts (of a process)
 - The program code, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time





Process in Memory

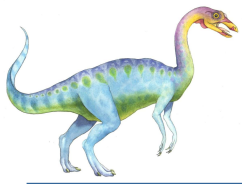




Process Concept (Cont.)

- ❑ A program is loaded and Execution of program started **via GUI mouse clicks, command line** entry of its name, etc
- ❑ One program can be several processes
 - ❑ Consider multiple users executing the same program (multiple copies)
 - ❑ Or the same user may invoke many copies of the web browser program
 - ❑ Each of these is a separate process; and
 - ❑ although the text sections are equivalent, the data, heap, and stack sections vary.
- ❑ It is also common to have **a process that spawns many processes** as it runs.





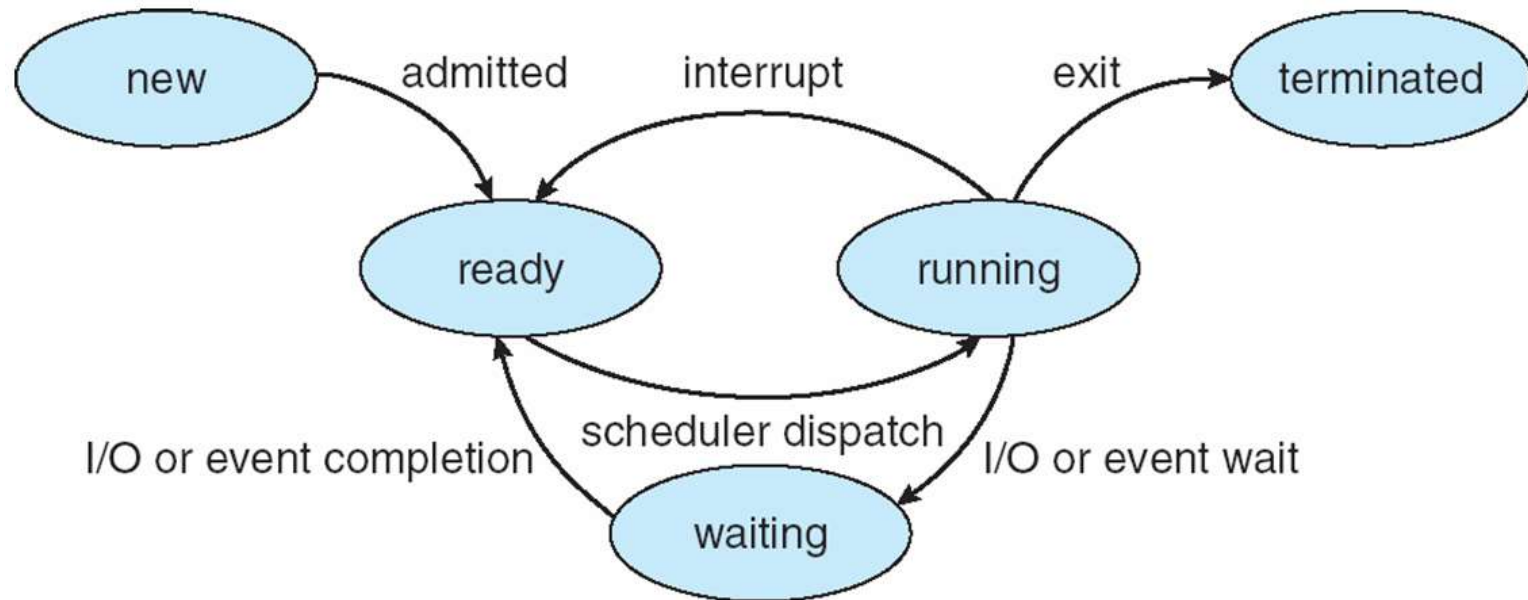
Process State

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution





Diagram of Process State

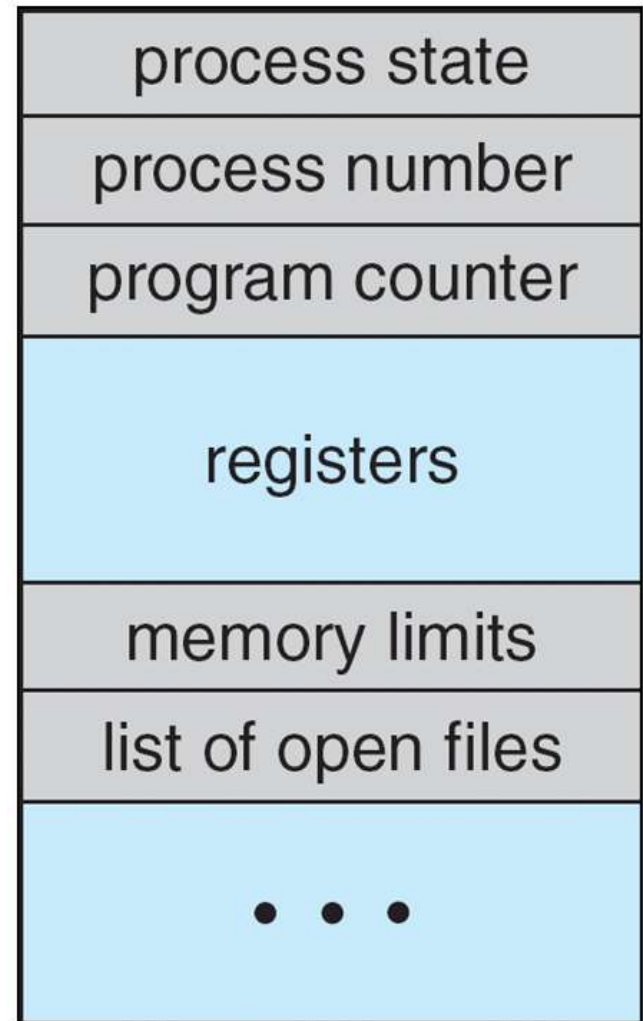




Process Control Block (PCB)

Information associated with each process
(also called **task control block**)

- ❑ Process state – running, waiting, etc
- ❑ Program counter – location of instruction to next execute
- ❑ CPU registers – contents of all process-centric registers
- ❑ CPU scheduling information- priorities, scheduling queue pointers ([chap 6](#))
- ❑ Memory-management information – memory allocated to the process ([chap 8](#))
- ❑ Accounting information – CPU used, clock time elapsed since start, time limits
- ❑ I/O status information – I/O devices allocated to process, list of open files





Threads

- ❑ So far, process has a single thread of execution
- ❑ The user cannot simultaneously type in characters and run the spell checker within the same process, for example.
- ❑ Modern operating systems extended the process concept to **allow a process to have multiple threads** of execution and thus to perform more than one task at a time.
- ❑ Consider having multiple program counters per process
 - ❑ Multiple locations can execute at once
 - ▶ Multiple threads of control -> **threads**
- ❑ Must then have storage for thread details, **multiple program counters in PCB**
- ❑ **Other changes are also needed** to support threads.
- ❑ See next chapter (chap 4)

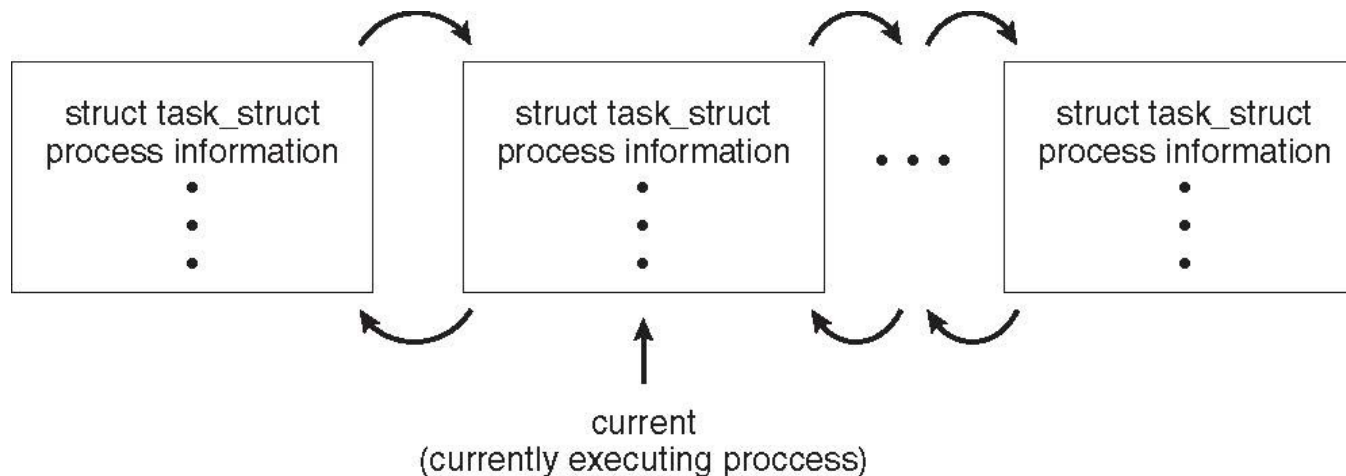




Process Representation in Linux

Represented by the C structure `task_struct`

```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```





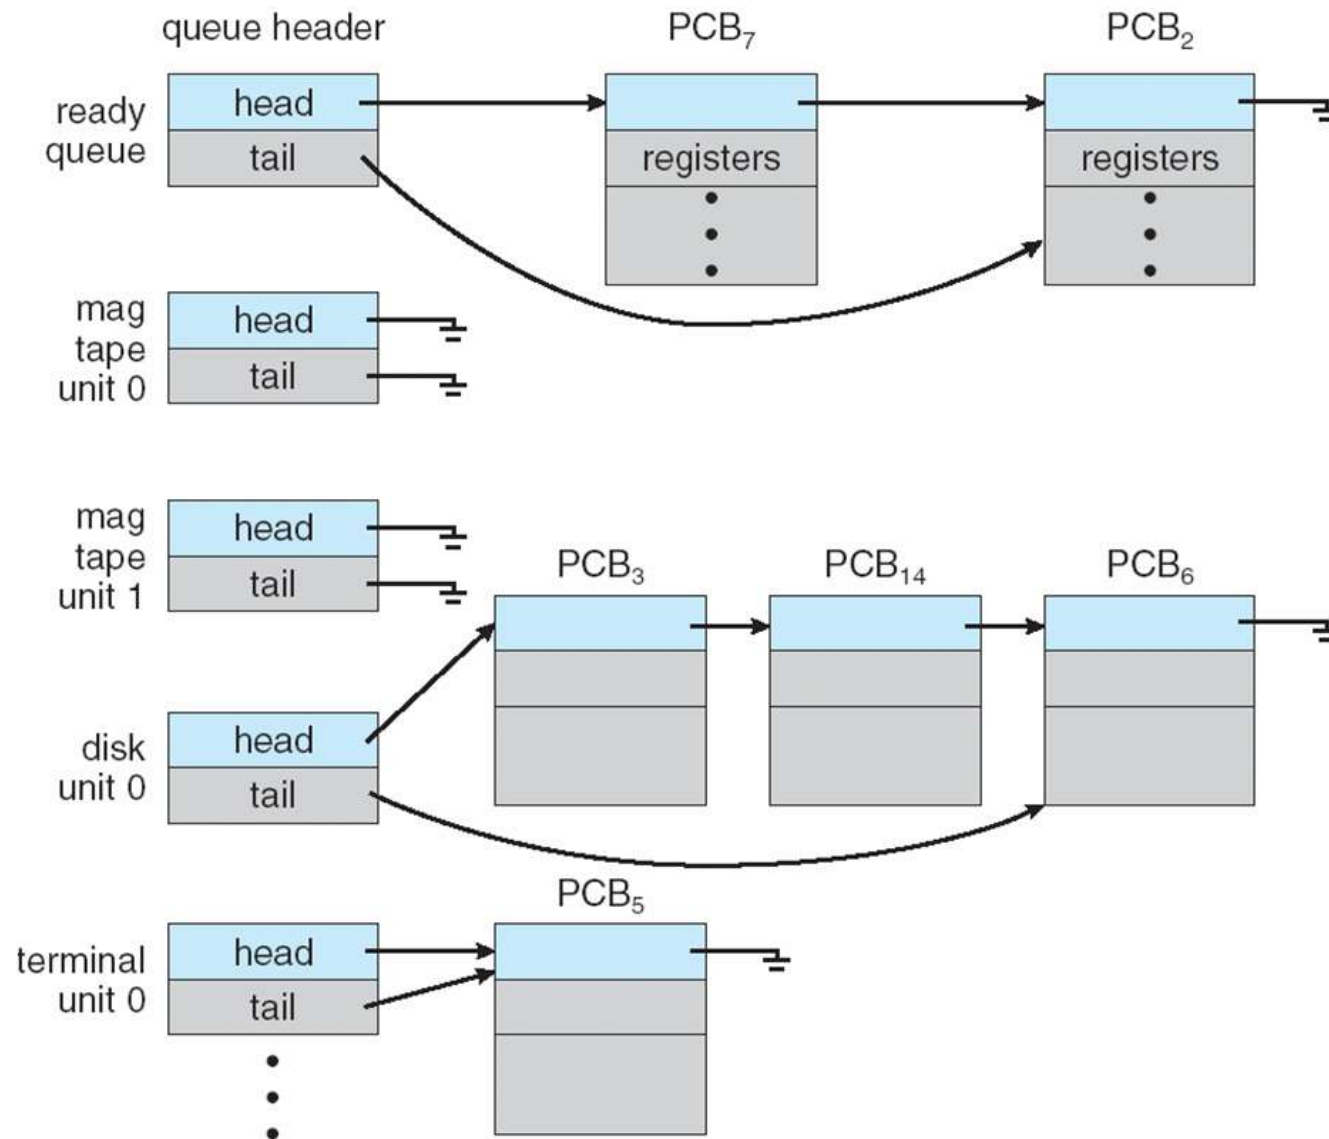
Process Scheduling

- ❑ Maximize CPU use, quickly switch processes onto CPU for time sharing
- ❑ **Process scheduler** selects among available processes for next execution on CPU
- ❑ Maintains **scheduling queues** of processes
 - ❑ **Job queue** – set of all processes in the system
 - ❑ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - ❑ **Device queues** – set of processes waiting for an I/O device (Each device has its own device queue)
 - ❑ **Processes migrate** among the various queues





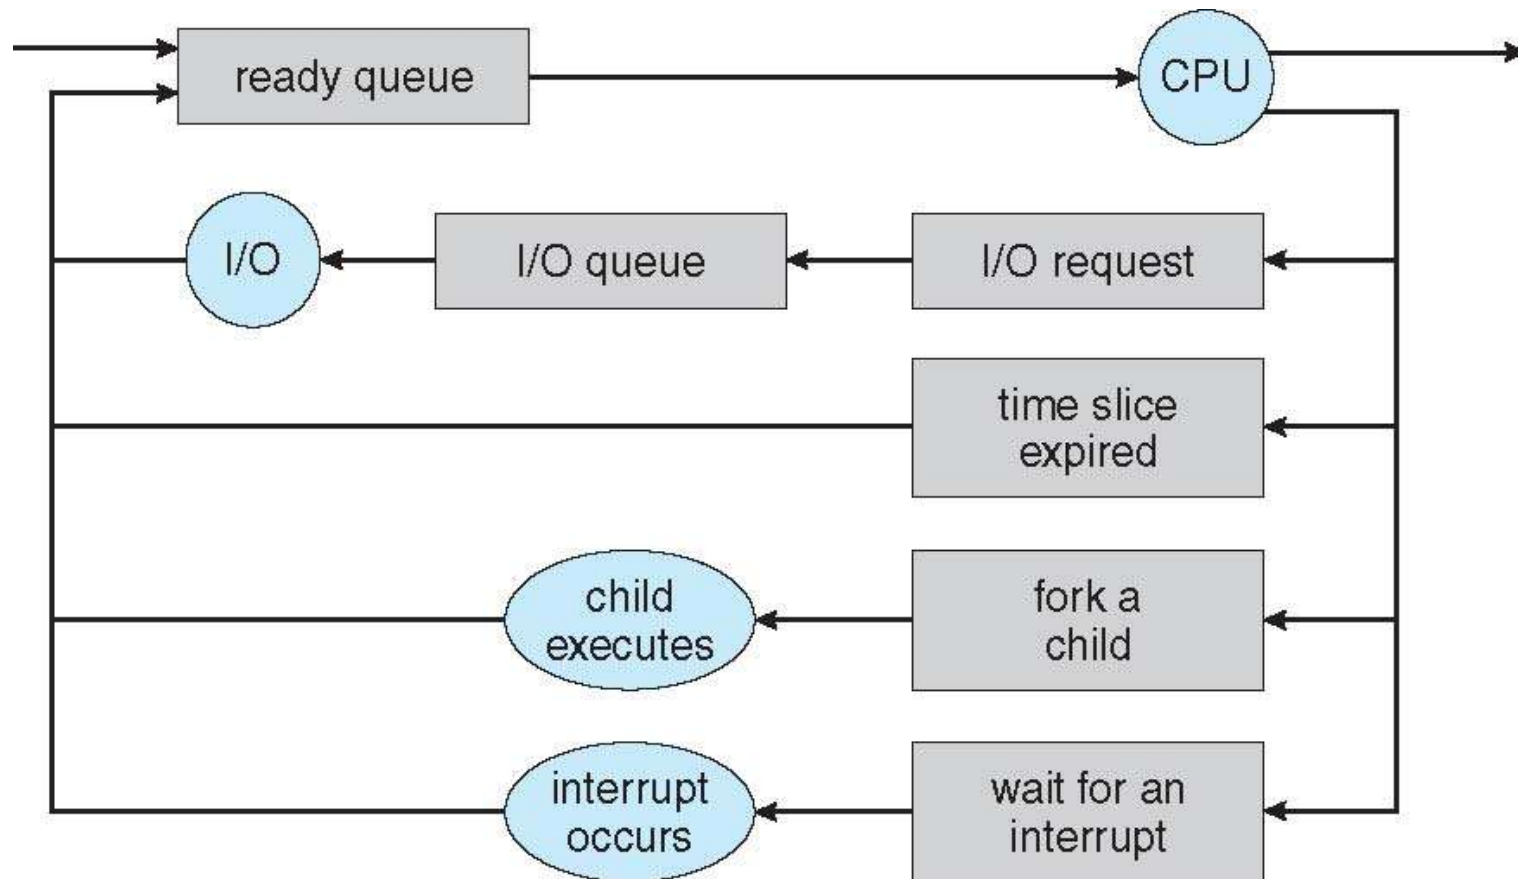
Ready Queue And Various I/O Device Queues





Representation of Process Scheduling

- **Queueing diagram** represents queues, resources, flows

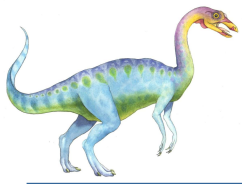




Schedulers

- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds) \Rightarrow (must be fast)
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes) \Rightarrow (may be slow)
 - The long-term scheduler controls the **degree of multiprogramming (the number of processes in memory)**
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good ***process mix***

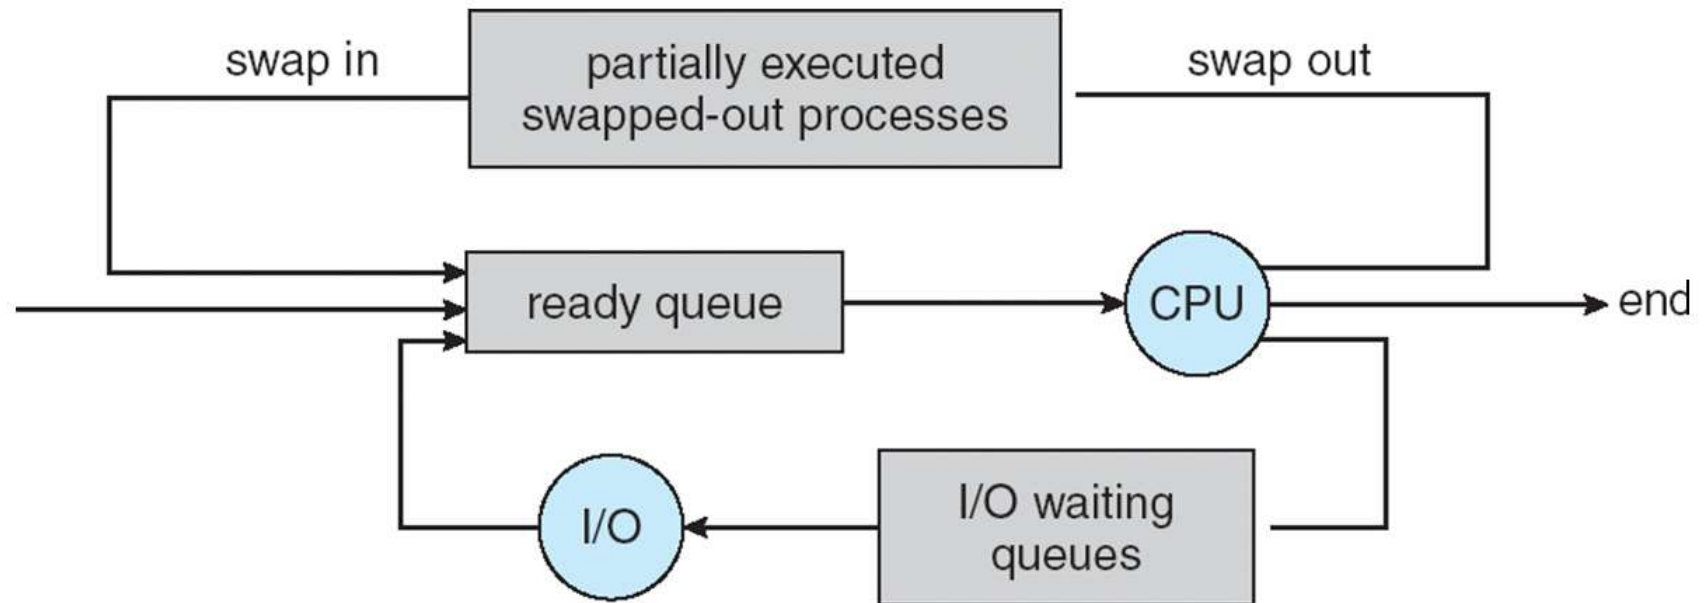


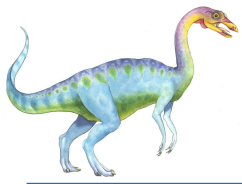


Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
- Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**
- Swapping may be necessary
 - to improve the process mix (of CPU bound and I/O bound)
 - or because a change in memory requirements has overcommitted available memory, requiring memory to be freed up.







Multitasking in Mobile Systems

- ❑ Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- ❑ Due to screen real estate, user interface limits iOS provides for a
 - ❑ Single **foreground** process- controlled via user interface
 - ❑ Multiple **background** processes– in memory, running, but not on the display, and with limits
 - ❑ Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
 - ❑ Apple probably limits multitasking due to battery life and memory use concerns.
- ❑ Android does not place such constraints on the types of applications that can run in the background
- ❑ Android runs foreground and background, with fewer limits
 - ❑ Background process uses a **service** to perform tasks (a separate application component that runs on behalf of the background process.)





Multitasking in Mobile Systems cont.

- Consider a streaming audio application:
- if the application moves to the background, the service continues to send audio files to the audio device driver on behalf of the background application.
- Service can keep running even if background process is suspended
- Service has no user interface, small memory use





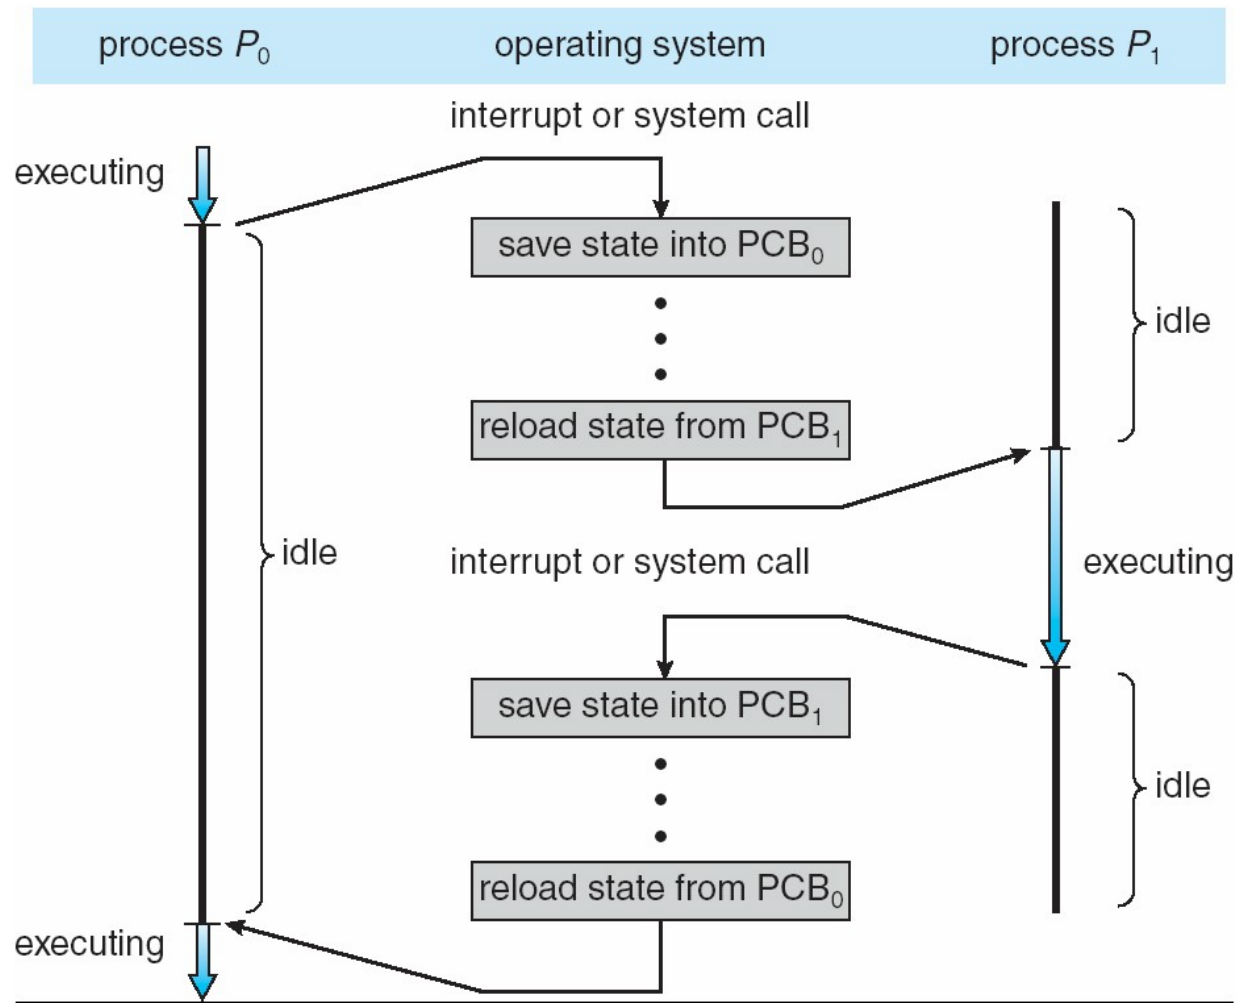
Context Switch

- When CPU switches to another process, the system must **save the state (in its PCB)** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support
 - Some hardware (e.g. Sun UltraSPARC) provides multiple sets of registers per CPU → multiple contexts loaded at once.
 - A context switch here simply requires changing the pointer to the current register set.





CPU Switch From Process to Process

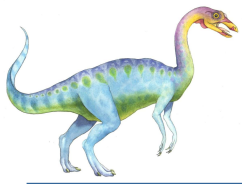




Operations on Processes

- System must provide mechanisms for:
 - process creation,
 - process termination,
 - and so on as detailed next





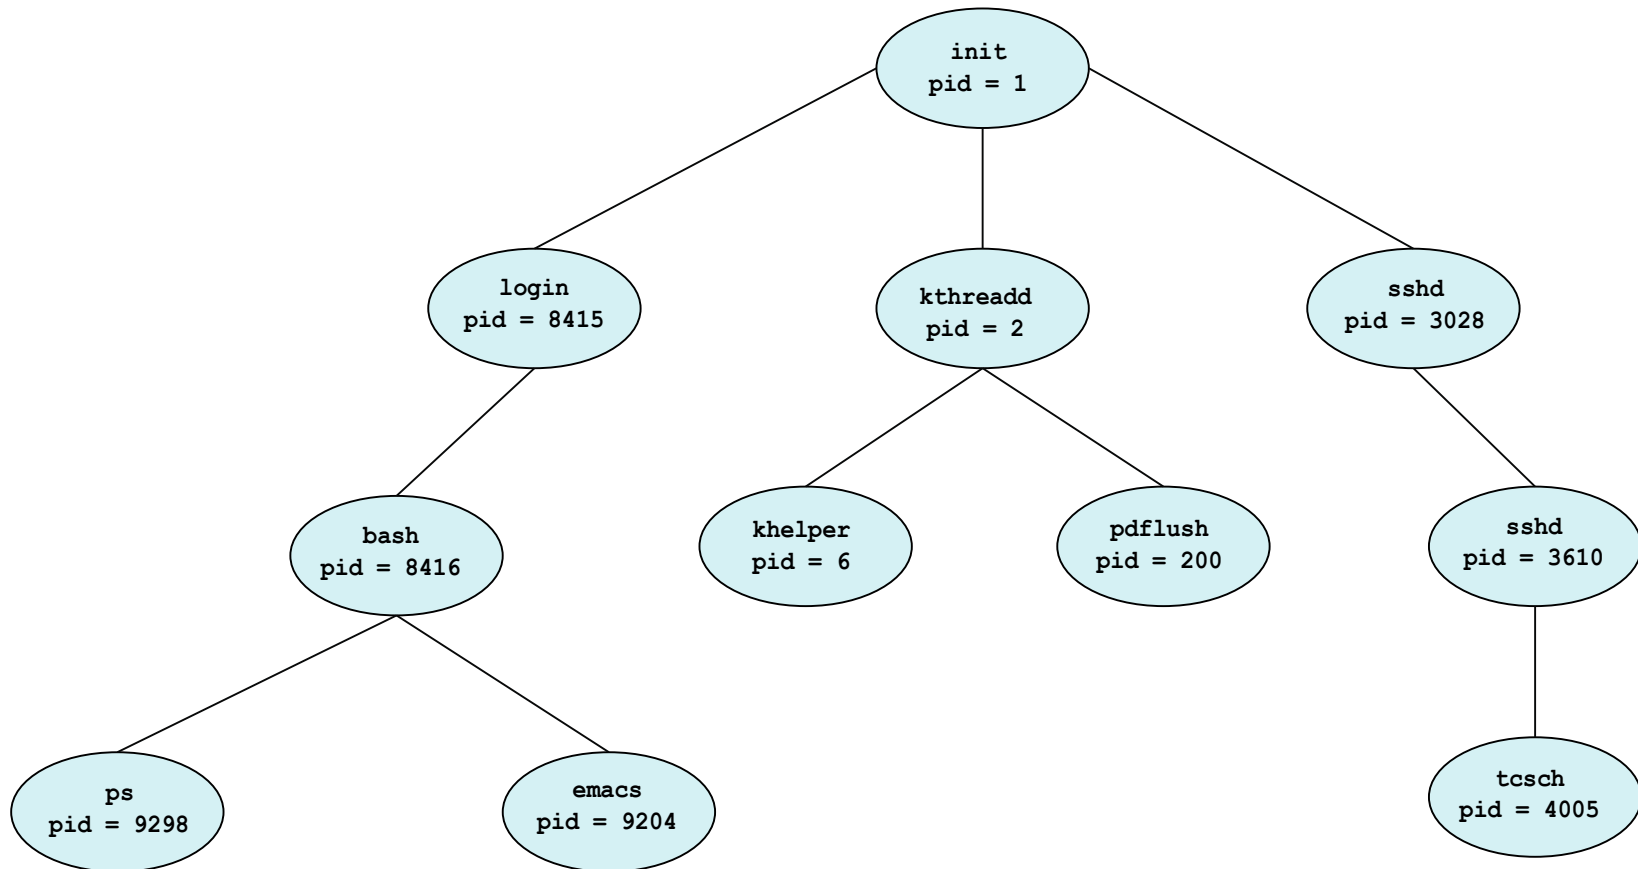
Process Creation

- ❑ **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- ❑ A child process will need certain resources (CPU time, memory, files, I/O devices)
- ❑ Generally, process identified and managed via a **process identifier (pid)**
- ❑ Resource sharing options
 - ❑ Parent and children share all resources
 - ❑ Children share subset of parent's resources
 - ❑ Parent and child share no resources (the child process obtains its resources directly from the operating system)
- ❑ Execution options
 - ❑ Parent and children execute concurrently
 - ❑ Parent waits until children terminate






A Tree of Processes in Linux





Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program





C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```





Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\\\WINDOWS\\\\system32\\\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```





Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
 - Returns status data from child to parent (via `wait()`)
 - Process' resources are deallocated by operating system
- **Parent may terminate** the execution of **children processes** using the `abort()` system call. Some reasons for doing so:
 - Child has **exceeded allocated resources**
 - Task assigned to child is **no longer required**
 - **The parent is exiting** and the operating systems does not allow a child to continue if its parent terminates





Process Termination

- ❑ Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - ❑ **cascading termination.** All children, grandchildren, etc. are terminated.
 - ❑ The termination is initiated by the operating system.
- ❑ The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

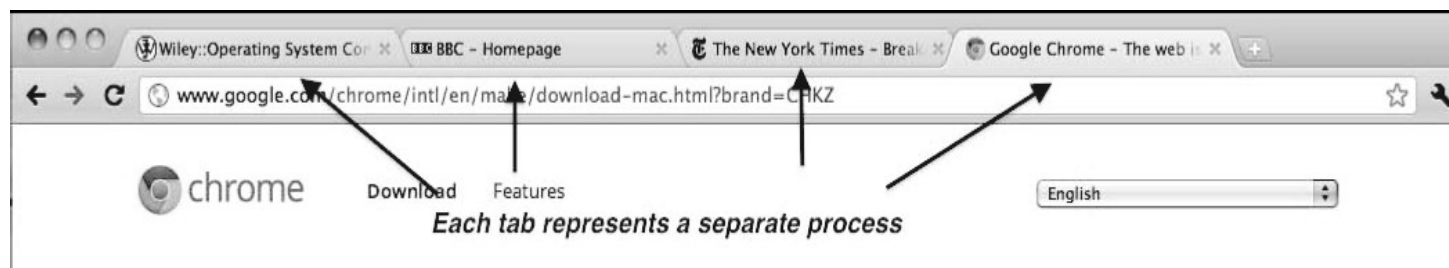
- ❑ If a process terminates but has no parent waiting (did not invoke `wait()`) process is a **zombie (because its entry in the process table entry continues to exist)**.
- ❑ If parent terminated without invoking `wait`, process is an **orphan (in Linux and Unix the init process becomes the parent which periodically invokes `wait()`, to allow the exit status of the orphan to be collected and thus to release the orphan's process table entry)**





Multiprocess Architecture – Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - **Browser** process manages user interface, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in





Interprocess Communication

- Processes within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing (e.g. a shared file between several users)
 - Computation speedup (to break a task into several tasks to run them n parallel)
 - Modularity (dividing the systems functions into separate processes)
 - Convenience (For instance, a user may be editing, listening to music, and compiling in parallel)





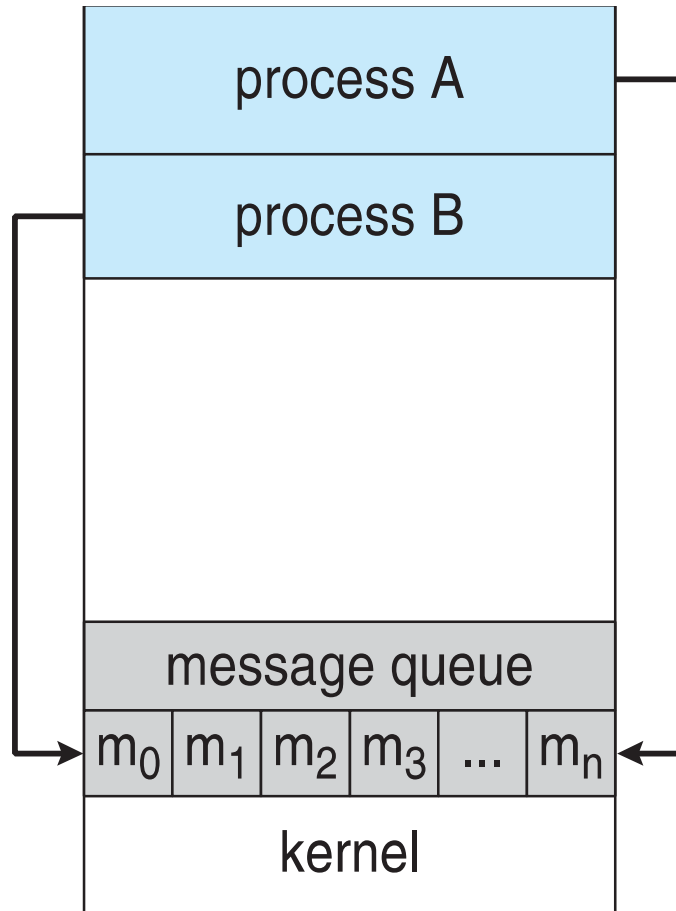
- ❑ Cooperating processes need **interprocess communication (IPC)** mechanism that will allow them to exchange data and information.
- ❑ Two models of IPC
 - ❑ **Shared memory**
 - ▶ A region of memory that is shared by cooperating processes is established.
 - ▶ Processes can then exchange information by reading and writing data to the shared region
 - ❑ **Message passing**
 - ▶ communication takes place by means of messages exchanged between the cooperating processes
 - ▶ Useful for exchanging smaller amounts of data, because no conflicts need be avoided and
 - ▶ easier to implement.
- ❑ Shared memory can be faster than message passing, since message-passing systems are typically implemented using system calls (time consuming)
- ❑ However recent research show that in systems with several cores message passing is faster (due to caches coherency issues)






Communications Models

(a) Message passing. (b) shared memory.



(a)



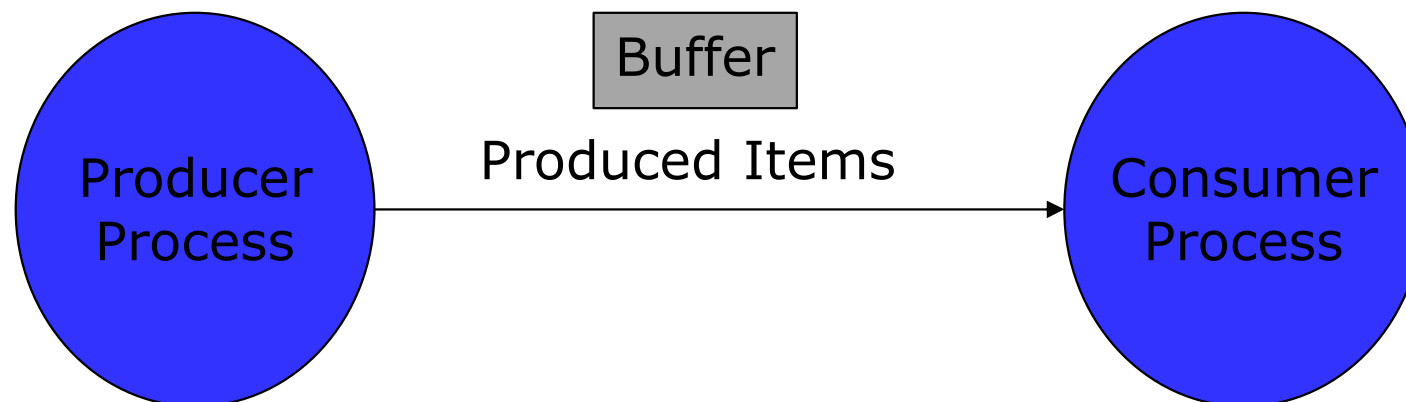
(b)





Producer-Consumer Problem

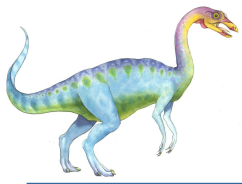
- To illustrate use of the **shared memory IPC mechanism**, a general model problem, called producer-consumer problem, can be used.
- *producer* process produces information that is consumed by a *consumer* process, e.g., a compiler produces code and the assembler consumes it





- Client-server metaphor
 - **unbounded-buffer** places no practical limit on the size of the buffer (the consumer may have to wait but the producer does not have to wait)
 - **bounded-buffer** assumes that there is a fixed buffer size (both may have to wait; the producer has to wait if the buffer is full, while the consumer has to wait if the buffer is empty)
- The shared buffer **is implemented as a circular array** with
 - two logical pointers: **in** and **out**.
 - The variable **in points to the next free position** in the buffer;
 - **out points to the first full position** in the buffer.
 - The buffer is empty when $in == out$;
 - the buffer is full when $((in + 1) \% BUFFER\ SIZE) == out$.





Bounded-Buffer – Shared-Memory Solution

□ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;


item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

□ Solution is correct, but can only use BUFFER_SIZE-1 elements






Buffer State in Shared Memory






Buffer State in Shared Memory

Buffer Full



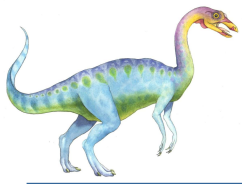
$$((in + 1) \% BUFFER_SIZE) == out$$

Buffer Empty



$$in == out$$





Bounded-Buffer – Producer

```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing - no free buffers */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```



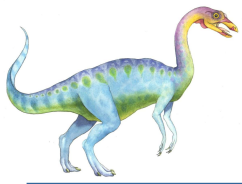


Bounded Buffer – Consumer

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing - nothing to consume */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next_consumed */
}
```






Interprocess Communication – Shared Memory

- ❑ An area of memory shared among the processes that wish to communicate
- ❑ The communication is **under the control of the users processes not the operating system.**
- ❑ Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory (i.e. how can they access the shared buffer concurrently; solution in chap 5)
- ❑ Synchronization is discussed in great details in Chapter 5.





Interprocess Communication – Message Passing

- ❑ Message passing provides a mechanism to allow processes to communicate and to synchronize their actions **without sharing the same address space.**
- ❑ It is particularly useful in **a distributed environment,**
- ❑ where the communicating processes may reside on different computers connected by a network.
- ❑ For example, **an Internet chat program**





Interprocess Communication – Message Passing

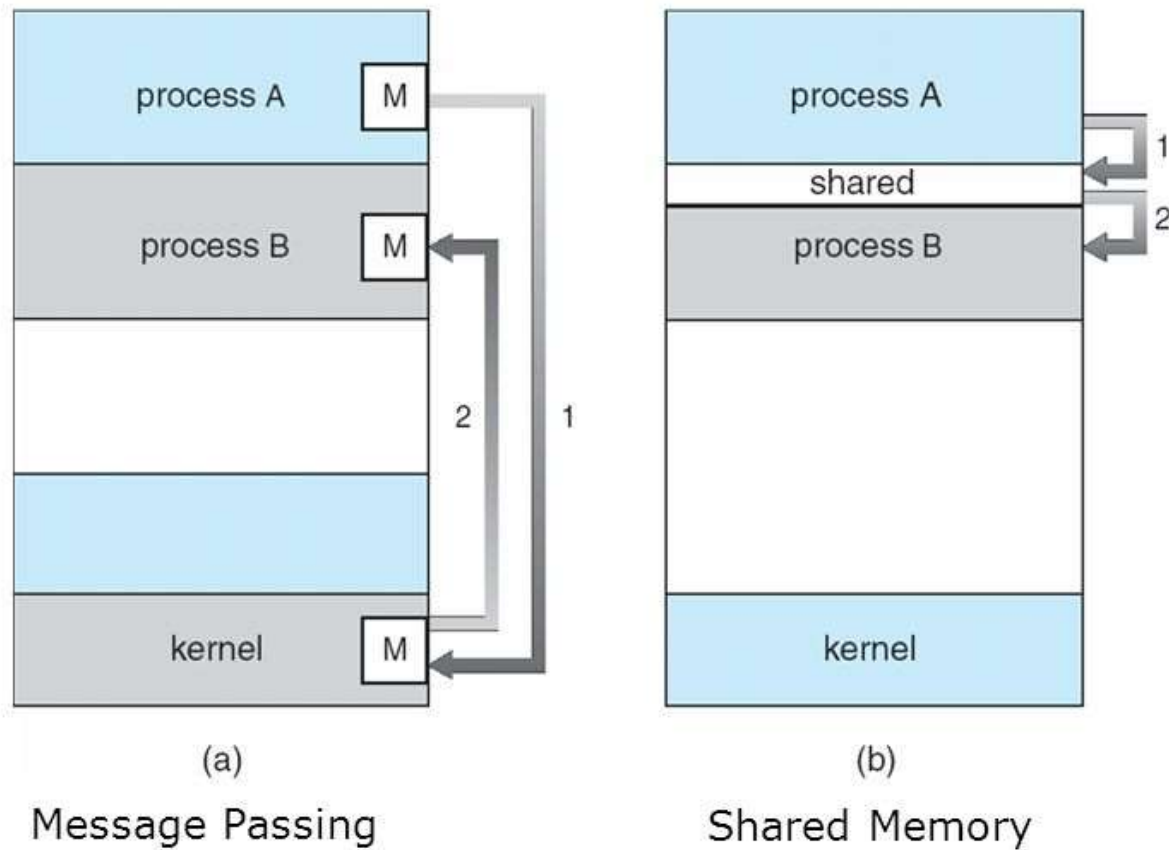
- ❑ Mechanism for processes to communicate and to synchronize their actions
- ❑ Message system – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations:
 - ❑ **send**(*message*)
 - ❑ **receive**(*message*)
- ❑ The *message* size is either fixed or variable
- ❑ Fixed-sized implementation is straightforward; however, it makes the task of programming more difficult
- ❑ Variable length messages require a more complex system level implementation, but the programming task becomes simpler
- ❑ A common tradeoff in OS design





Interprocess Communication – Summary

Communications Models



End of Chapter 3

