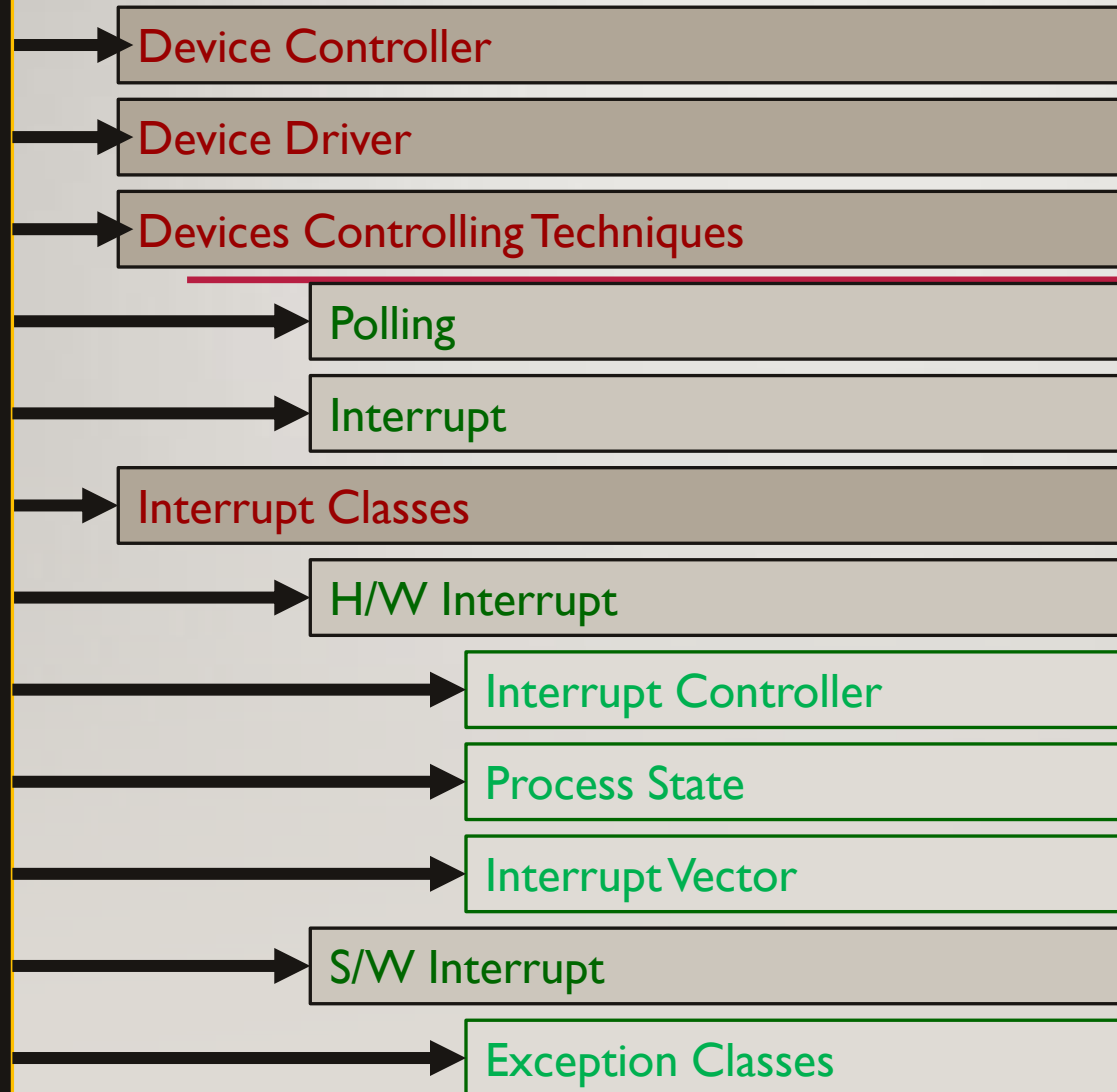


CHAPTER I

---

# INTERRUPTS




## DEVICE CONTROLLER


Each I/O unit consists of two components: a **mechanical** one and an **electronic** one.

The electronic one is called the **device controller** or **device adapter**.

On personal computers, it often takes the form of a **printed circuit card** that can be inserted into an **expansion slot** on the mainboard.



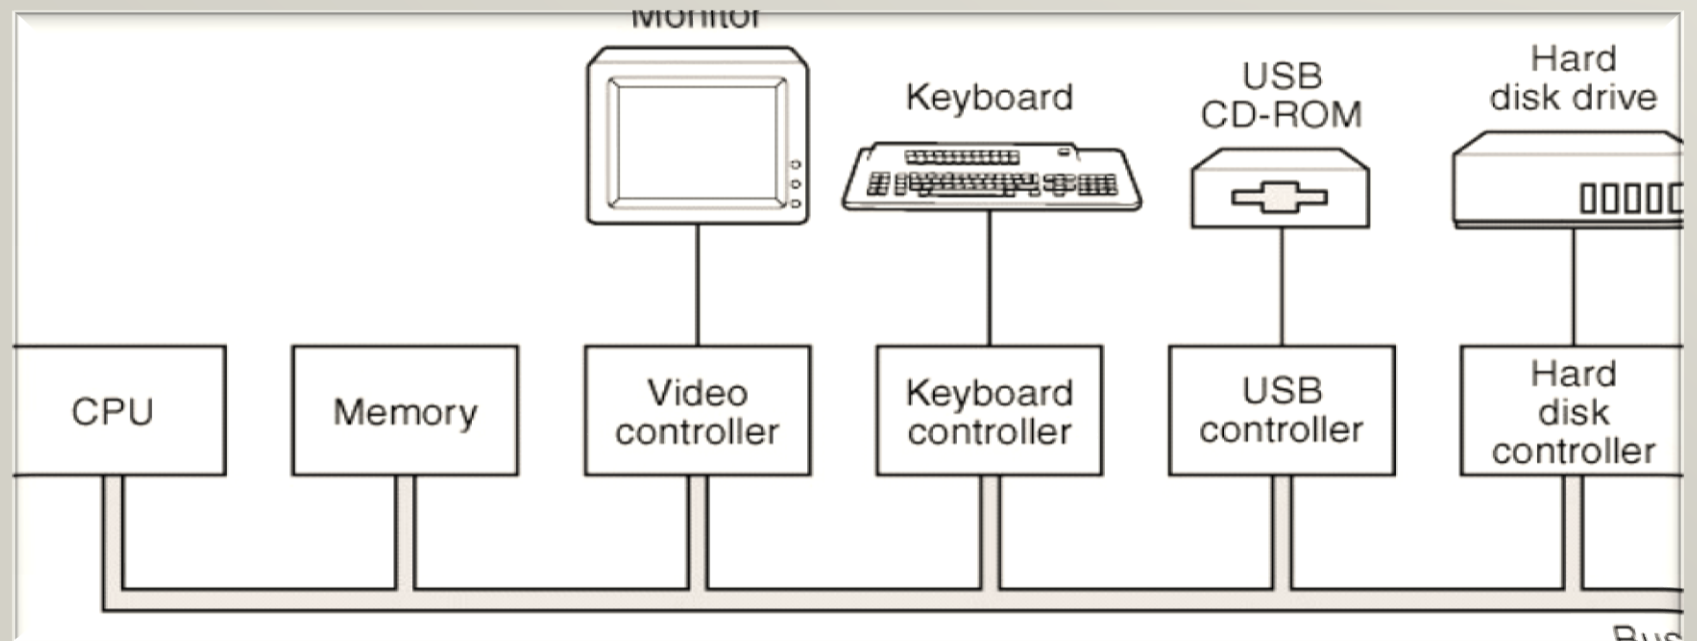
Device Controller (Adapter)



Motherboard (mainboard)

## DEVICE DRIVER

The controller card usually has a **connector** into which a cable leading to the device itself can be plugged.



The OS deals with device controller (H/W) rather than the device itself through the corresponding **device driver** (S/W).

Each controller has a few registers. Based on the contents of these registers, the OS can command the device to perform a specific action such as deliver data, accept data, switch on/off, etc...

Also, the OS learns the state of the device from these registers.

In addition to the registers, a controller contains data buffers that the OS can read/write.

## DEVICES CONTROLLING TECHNIQUES

It is important for the OS to **recognize and manipulate** the status of each device.

The status of a device includes the following information:

- Is it correctly connected or not?
- Does it work or not?
- Is it busy or available?

The OS keeps track of the statuses of all connected devices using a device-status table stored in memory.

Two techniques are applied for this purpose:

- Polling
- Interrupt

## POLLING

### ~ Target:

In order for the CPU to recognize the status of each connected device, a communication should take place between both.

### ~ Concept:


The CPU repeatedly requests from each device its current status.

### ~ Definition:

The CPU sends electrical signals regularly to each connected device asking about its current status. The device emits another signal to respond to each signal sent by the CPU.

### ~ Disadvantages:

- \* This causes an overhead on the CPU
- \* The CPU stays idle → waste of resources
- \* The mechanism is slow
- \* Increases the complexity of the computer system



**Interrupts** eliminate the need for a processor to repeatedly poll devices.

## INTERRUPTS

### ~ Definition:

An interrupt is an electrical signal emitted by the hardware or software to the processor indicating an event that needs immediate attention.

### ~ Target:

Interrupts allow the OS to take control of the processor to manage system resources.

### ~ Concept:

An interrupt alerts the processor to a high-priority condition requiring the “interruption” of the current code that the CPU is executing.

As soon as the CPU receives an interrupt, the following steps are performed:

- (1) The processor suspends its current activity.
- (2) The processor saves its **state**.
- (3) The processor calls the **Interrupt Handler (IH)**, sometimes called **Interrupt Service Routine (ISR)**, to handle the event.

The **Interrupt Handler (IH)** – or **Interrupt Service Routine (ISR)** – is a relatively small program (procedure). It is part of the OS.

There are many types of interrupts. Each type of interrupt calls the corresponding Interrupt Handler.

## INTERRUPT CLASSES

```
graph TD; A[INTERRUPT CLASSES] --> B[Hardware Interrupts]; A --> C[Software Interrupts];
```

Hardware Interrupts

Software Interrupts



## INTERRUPT CLASSES – H/W INTERRUPTS

- ~ Emitted by a device to indicate that it requires attention from the OS.
- ~ Implemented using electronic signals sent from an internal/external device.
- ~ For example, pressing a key on the keyboard “triggers” an interrupt that causes the processor to read the keystroke.
- ~ May occur in the middle of instruction execution.
- ~ **Asynchronous**: changes immediately – does not wait for the fall/rise of a clock cycle.
- ~ Requires special care in programming.
- ~ The act of initiating a hardware interrupt is known as an **Interrupt Request (IRQ)**.
- ~ The number of H/W interrupts is limited by the number of IRQ lines.

## H/W INTERRUPTS – INTERRUPT CONTROLLER

An **interrupt line** is an electrical connection between the mainboard and the CPU.

Each device has a **controller** that is responsible to activate the interrupt line when necessary (ie. when a status changes) to inform the CPU that an event has occurred.

The CPU is provided with an **interrupt controller** that sorts the received requests based on their priority.

Important interrupts are therefore serviced first. The rest of interrupts are queued and serviced after higher-priority ones.

## HARDWARE INTERRUPTS – PROCESS STATE

The paused process (program) resumes after the CPU handles the current interrupt.

In order to continue a process after being paused, some information should be saved.

Examples of such information include:

- ~ where did the execution stop (instruction address)
- ~ the values of data generated during process execution: these are saved in the registers.
- ~ and others...

These information are known as the **process state**.

The process state is stored in the **Process Control Block (PCB)**. Each process has its own (PCB).

## H/W INTERRUPTS – INTERRUPT VECTOR

Each type of interrupt is assigned a unique value that the CPU uses as an index into the **interrupt vector**.

The **interrupt vector** is an array of pointers to the memory addresses of all available interrupt handlers.

The interrupt vector is loaded in a protected area of the memory, so that it cannot be erased or modified inadvertently.

## H/W INTERRUPTS – MANIPULATION

### Step 1:

When the status of a device changes, it activates the interrupt line to inform the CPU about the event.

### Step 2:

As soon as the CPU receives the signal through the interrupt line, it completes execution of the current instruction then pauses the execution of the current process (program).

### Step 3:

The CPU saves the state of the paused process.

### Step 4:

The CPU transfers control to the appropriate interrupt handler.

### Step 5:

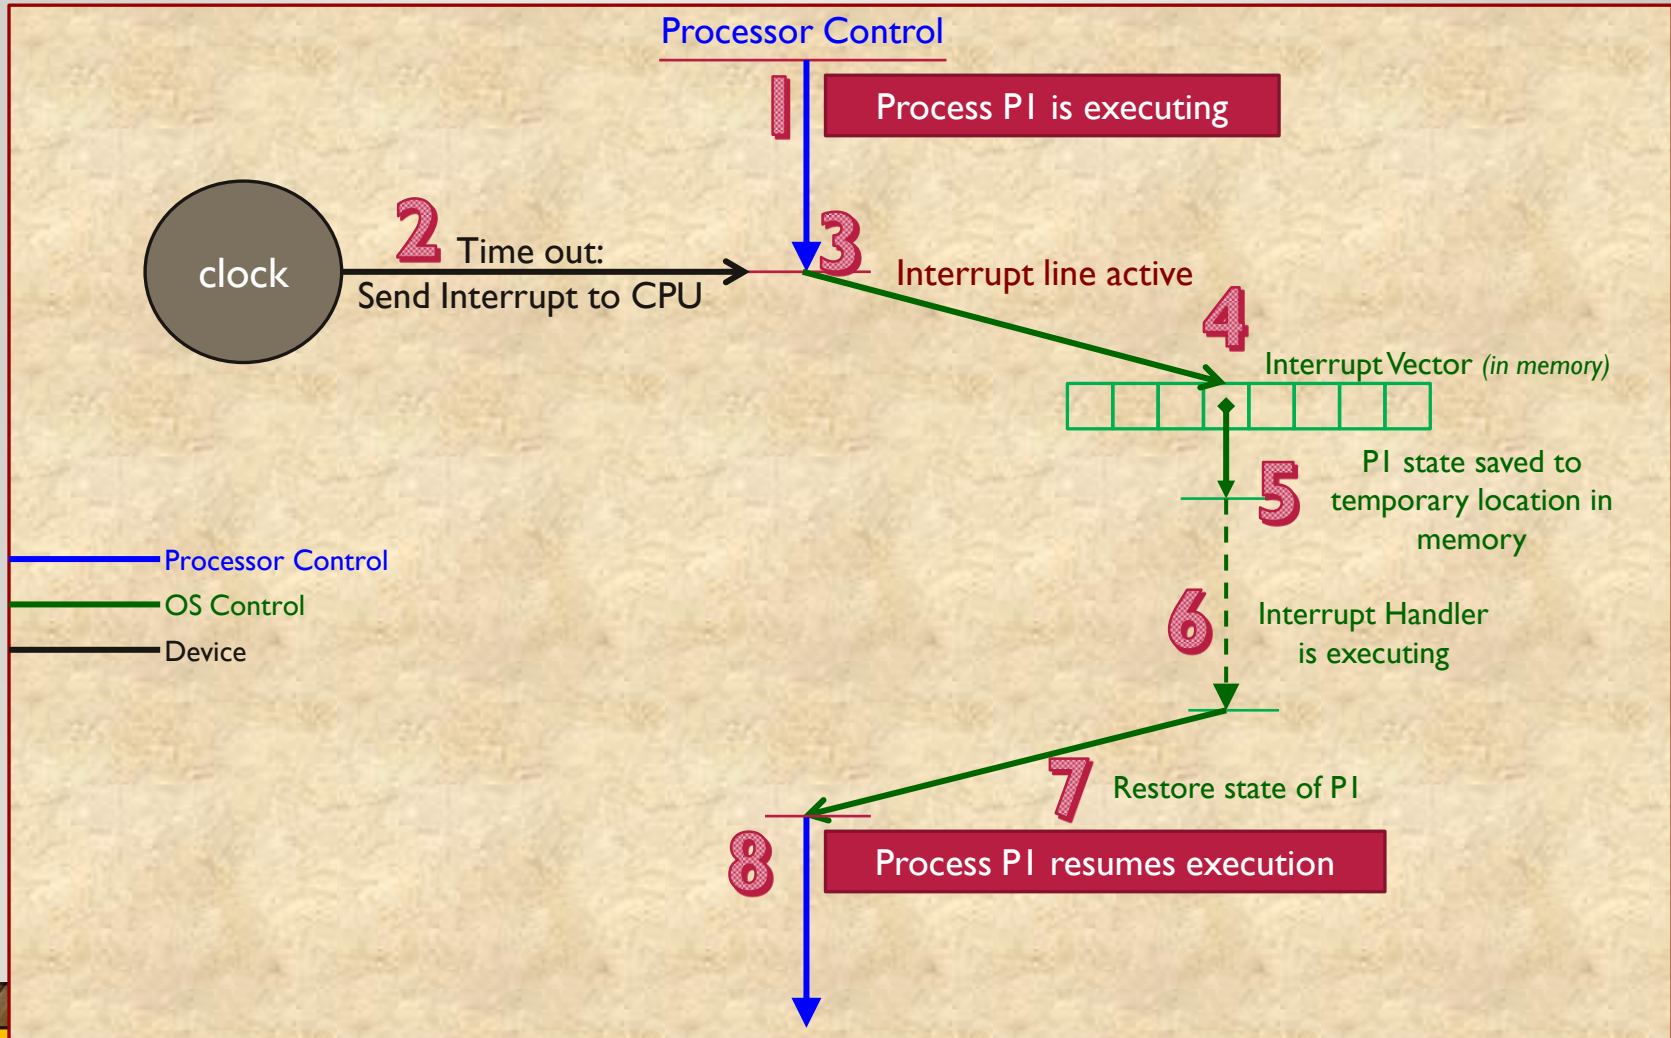
The interrupt handler performs appropriate actions based on the type of interrupt.

### Step 6:

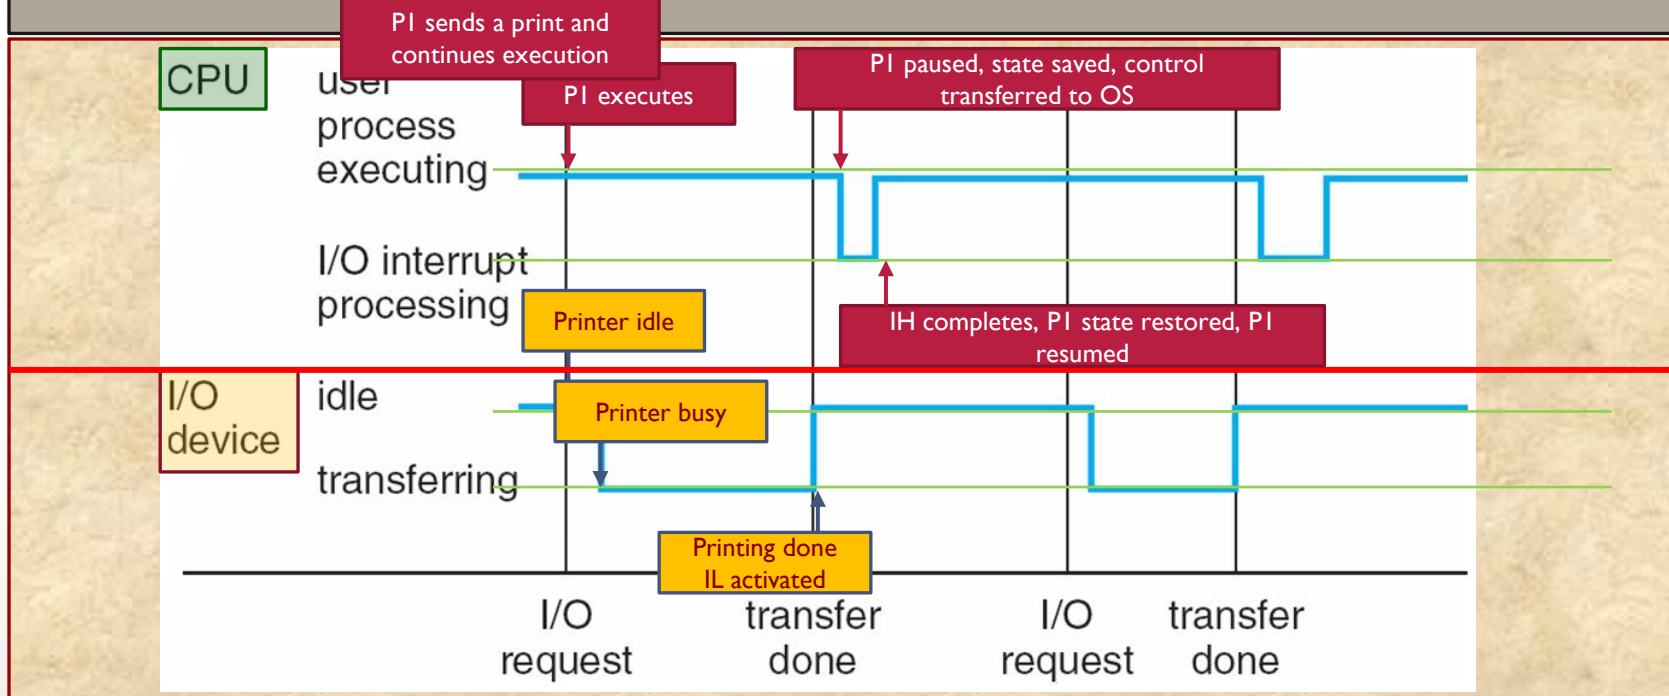
After the interrupt handler completes execution, the state of the previously interrupted process is restored.

## H/W INTERRUPTS – EXAMPLE

A time-out is set to a specific time. When the time elapses the clock sends an interrupt.



## H/W INTERRUPTS – INTERRUPT TIMELINE



IL = Interrupt Line

IH = Interrupt Handler

Note how the H/W-generated interrupts are asynchronous; i.e. Changes do not wait for the rise/fall of the clock.

## INTERRUPT CLASSES – S/W INTERRUPTS

~ May occur **unintentionally** due to an error that cannot be handled by the CPU: this is known as **exception** or **trap**.

~ For example, an exception is “**thrown**” when the program attempts to:

- \* Divide by zero.
- \* Access a protected memory area.
- \* Access an out-of-memory address.

~ In such cases, the OS decides to **abort** (end) the program.

~ S/W Interrupts may also be **intentionally** caused by a special instruction in the instruction set.

~ Examples of intentionally caused S/W interrupts include:

- \* Breakpoints set within a program for debugging purposes.
- \* Exceptions issued in Java code.

~ An OS may include up to hundreds of S/W interrupts for different purposes.



## S/W INTERRUPTS – EXCEPTIONS CLASSES

The set of S/W interrupts that a computer supports depends on the system's architecture.

For example, Intel IA-32 architecture supports a set of three exceptions:

- \* Faults
- \* Traps
- \* Aborts

The class of exception specifies the behavior of the running process.

**Faults** cause the running process to stop execution (abort). After the problem is corrected, the stopped process restarts.

**Traps** continue executing the running process after giving a message. Examples include breakpoints and overflow (the value stored in a register exceeds its capacity).

**Aborts** lead to a system crash. The OS may not be even able to store the state of the running process. Examples include hardware failure; or an exception handling itself causes another exception: this is called **double-fault exception**.

# H/W vs. S/W INTERRUPTS

## H/W-generated

- ~ Generated by H/W devices
- ~ Issued to inform the CPU that an event occurred (ie. H/W status changed)
- ~ Examples:
  - \* A keyboard issues an interrupt when the user presses a key
  - \* A printer issues an interrupt upon completion
- ~ Asynchronous:  
time-independent of running instruction

## S/W-generated

- ~ Generated by programs
- ~ Issued when an error occurs or intentionally for debugging purposes
- ~ Examples:
  - \* Division by zero
  - \* A H/W device is malfunctioning
  - \* Breakpoint/Exception in the code
  - \* Accessing a protected memory area (reserved for the OS)
  - \* Accessing an out-of-memory address
- ~ Synchronous:  
time-dependent of the running instruction