

CHAPTER 3

THE PROCESS CONCEPT

gettyimages®

Daniel Grill



PROCESS DEFINITION (I)

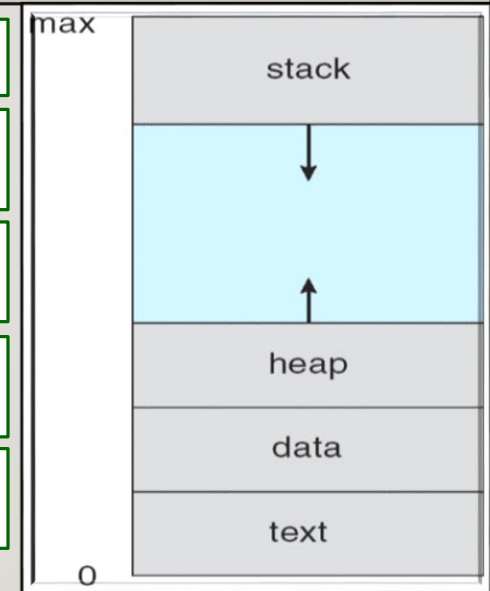
A **process** is a **program** in execution. This is also called a **job**.

A program is a passive entity; whereas the process is an active one.

In other words, a program becomes a process when it is loaded into memory.

A process includes all the information necessary to make it run during its lifetime such as:

- The **program code**, also known as the **text section**.
- The **program counter** that includes the address of the next instruction to be executed.
- The final contents of the **processor's registers** during the execution of the process.
- The **stack**: this contains the functions (methods) parameters, return addresses, and global variables.
- The **heap**: this contains the dynamically allocated memory required during run time.



PROCESS DEFINITION (2)

Although two processes may be associated with the same program code (text section), they are nevertheless considered as two separate execution sequences.

The following are some illustrating examples:

- Several users may be running different copies (processes) of the mail program.
- The same user may be running many copies (processes) of the web browser program.

In the above examples, the text sections are equivalent; but the data, heap, and stack versions vary.

A process itself may be an execution environment for other code:

As an example:

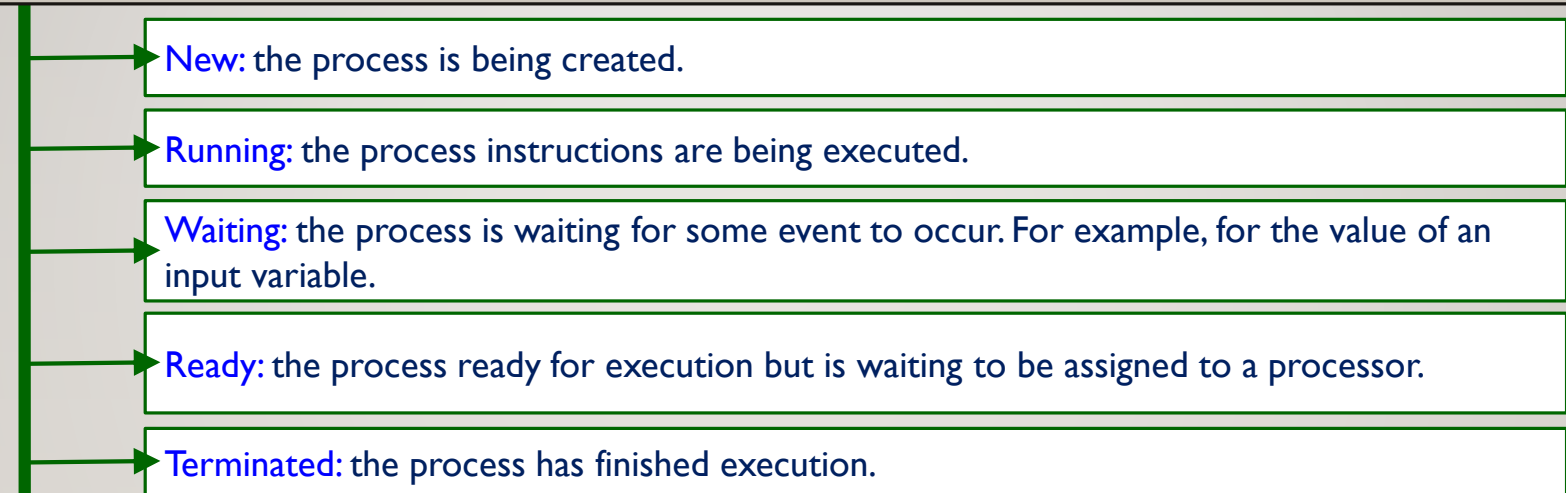
- The Java Virtual Machine (JVM) runs as a process.
- However, JVM interprets the loaded code (your program), and takes actions on behalf of that code.

PROCESS STATE (I) – INTRODUCTION

The **current** activity of a process is known as the **process state**.

This means that the process state continuously changes with time during a process lifetime.

A process may be in one of the following states during its lifetime:

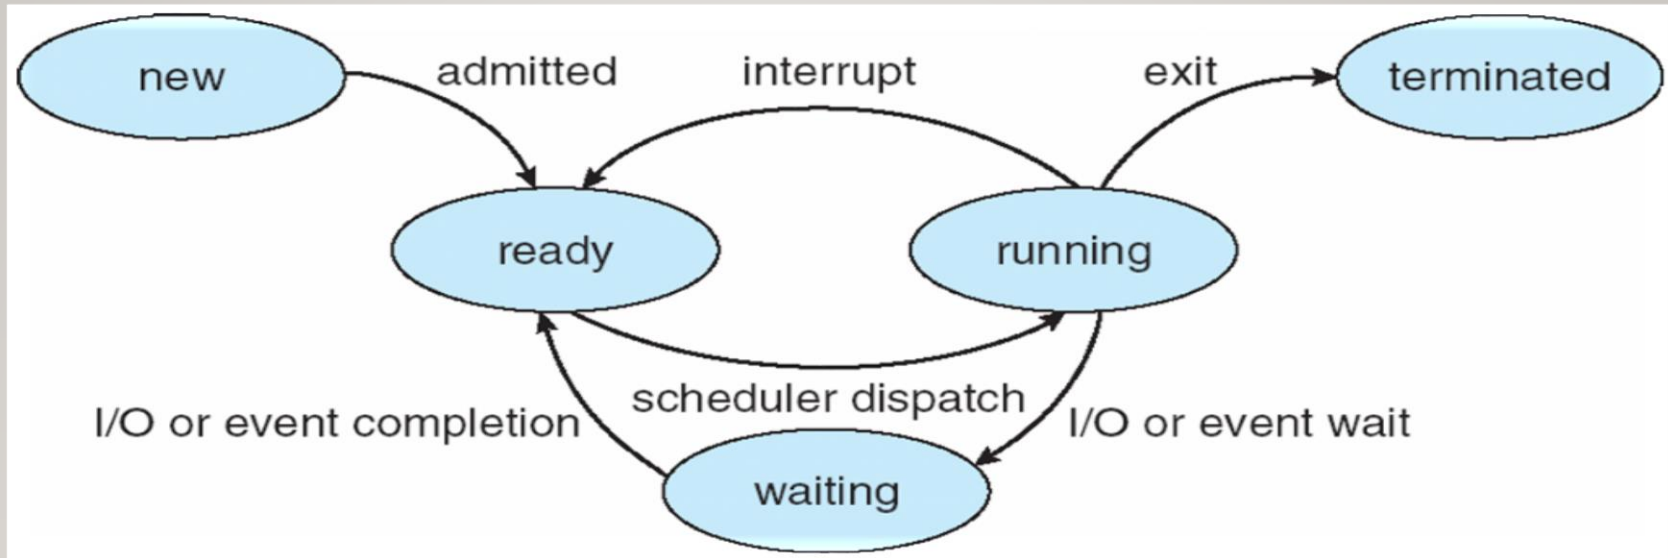


A processor executes at most a **single instruction** at any instant of time.

The instructions executed by a processor during a specific period of time might belong to one or more processes.

Note that many processes might be ready at an instant of time. In such case, they contend (compete) for the processor(s).

PROCESS STATE (2) – THE STATE DIAGRAM (I)



A state diagram consists of **states** and **actions**.

The process state diagram depicts the following:

→ All possible **states** that a process can go through during its lifetime.

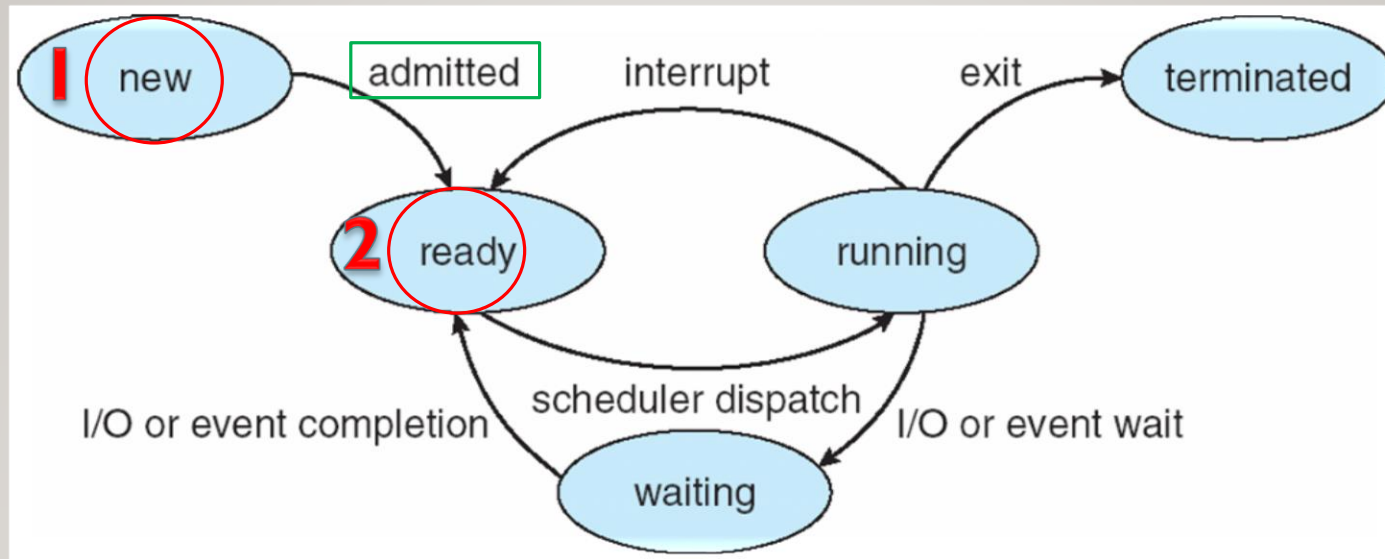
States are shown in **circles**.

→ All possible **actions** that cause a process to change its state.

Actions are shown with **arrows**.

→ A **transition** consists of two states: the previous state and the current one, plus the action that caused such state change.

For example, consider the following transition: a process in the running state changes to the **ready** state when **interrupted**.

PROCESS STATE (3) – THE STATE DIAGRAM (2) – THE NEW & READY STATES

A written program is initially on the disk in the **job pool**.

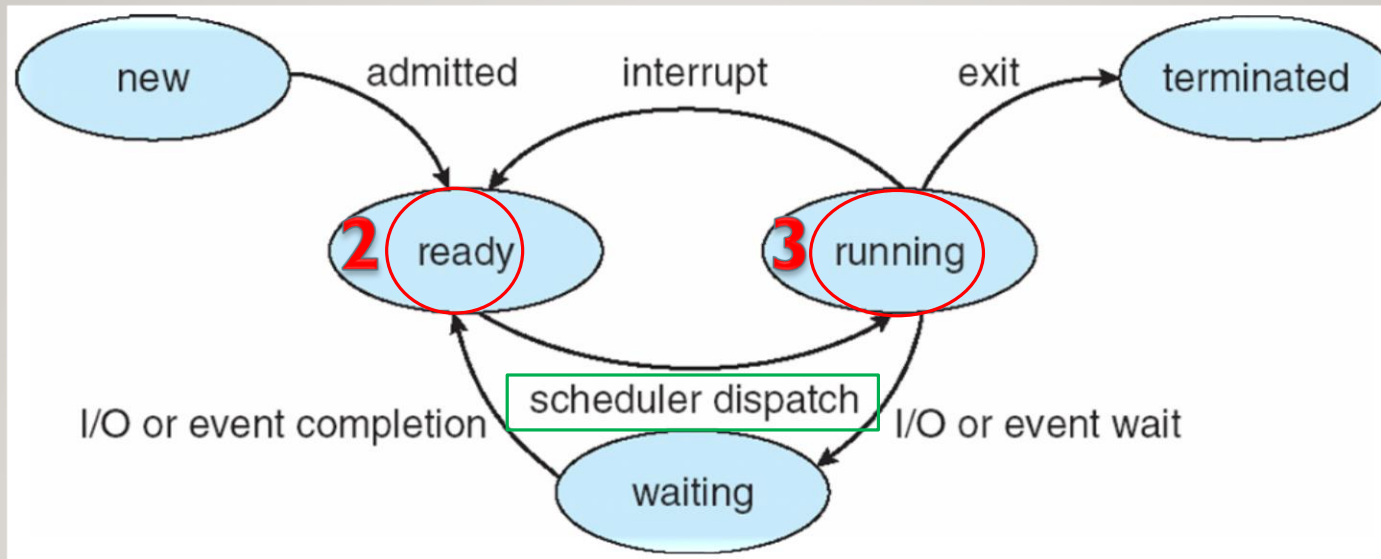
When the OS selects the program to be loaded into memory, a process is created.

At this instant of time, the process is given the status **new**.

When the OS creates a new process, it assigns to it a unique identifier (number). This is known as the **process identifier**.

The process is soon **admitted** and it becomes **ready** for execution.

PROCESS STATE (4) – THE STATE DIAGRAM (3) – THE READY & RUNNING STATES



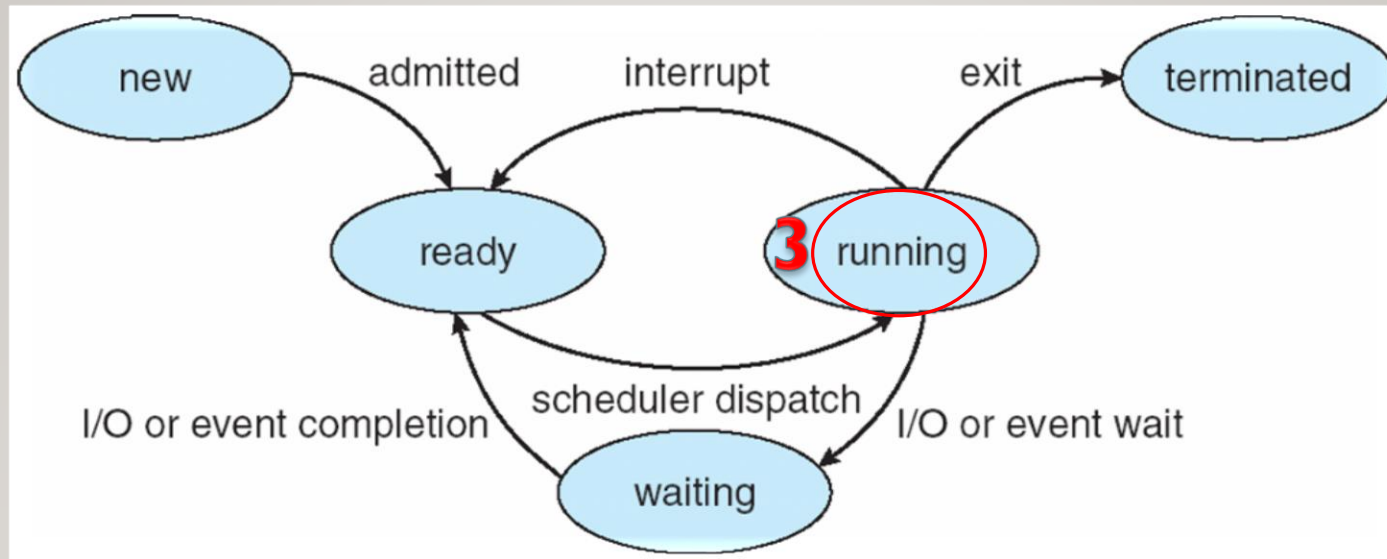
Therefore, a process in the **ready** state resides in the memory. We say that the process is in the **ready queue**. It waits the OS to assign him a processor to execute on.

Note that we may have multiple processes in the ready queue competing for the processor.

The OS selects only one process to be assigned the processor. Selection is made by a specific part of the OS called the **CPU scheduler**.

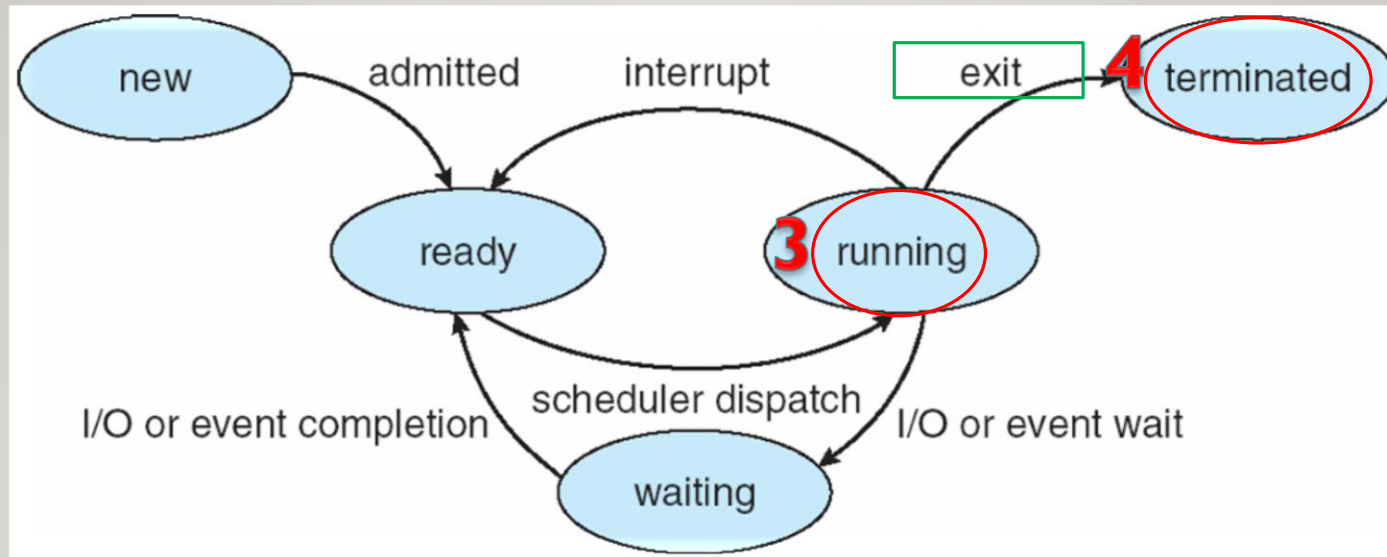
The CPU scheduler is made to select a process according to specific criteria.

As soon as the process is assigned a processor, it executes and changes to the **Running** state. We say that the **scheduler dispatches** the process.

PROCESS STATE (5) – THE STATE DIAGRAM (4) – THE RUNNING STATE

While running, a process may change its state to one of the following:

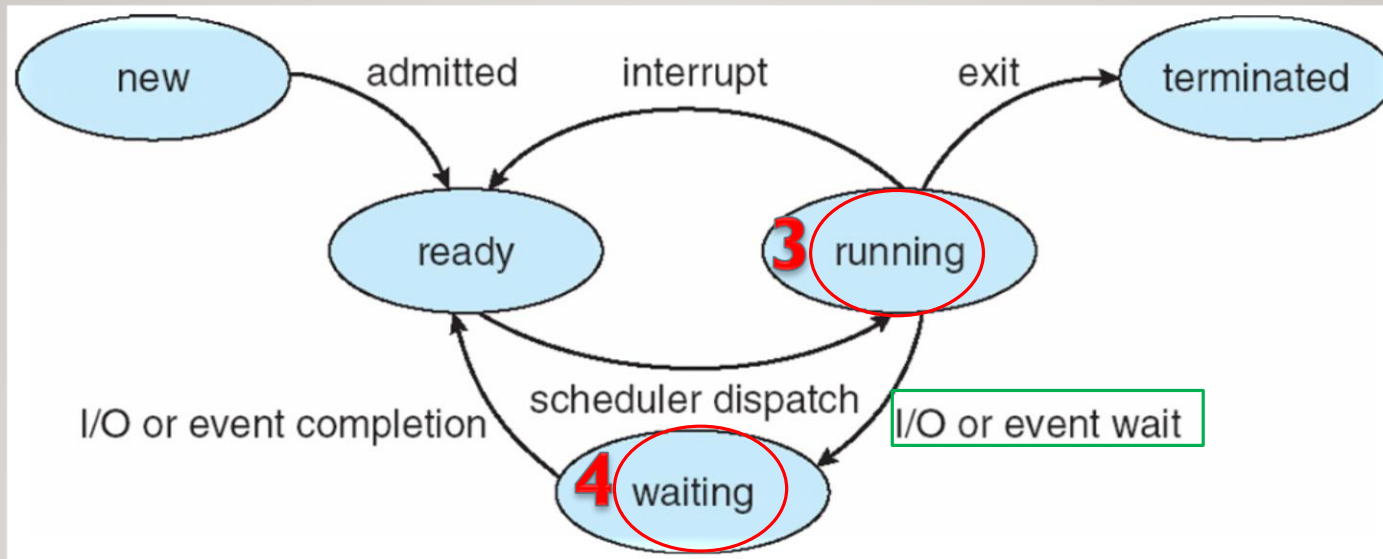
- Terminate
- Wait
- Ready (again...)

PROCESS STATE (6) – THE STATE DIAGRAM (5) – THE RUNNING & TERMINATED STATES

A process may run on the processor until it completes execution.

In this case, the process **exits**, changes to the **terminated** state.

A terminated process deallocates (leaves) the memory, and all previously allocated resources.

PROCESS STATE (7) – THE STATE DIAGRAM (6) – THE RUNNING & WAITING STATES

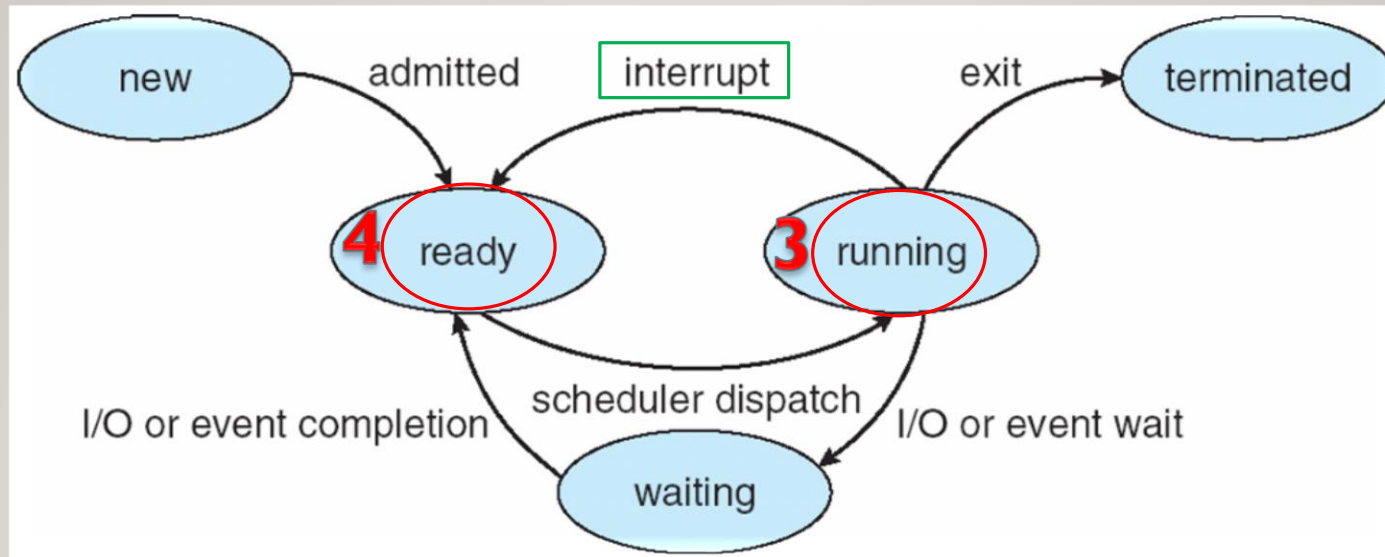
While running a process, a processor may encounter an I/O statement.

Assume for example, that the processor encountered an “input” statement.

In this case, the process halts execution until the user enters the input. The processor speed is millions times faster than pressing a single key on the keyboard.

In order to make full use of the processor, the process is appended to a device waiting queue. The process status then changes to the waiting state.

During the time the process is in the device waiting queue, the processor runs another ready process in the ready queue, if any. Otherwise, it stays idle.

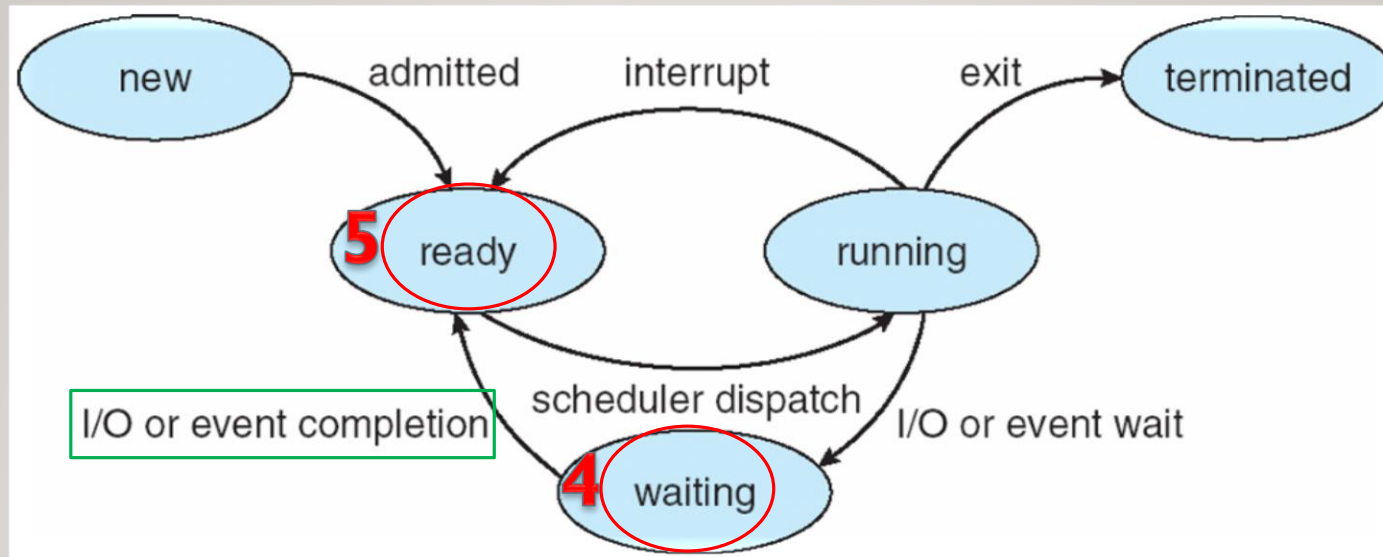
PROCESS STATE (7) – THE STATE DIAGRAM (6) – THE RUNNING & READY STATES

In a multi-process environment, the OS gives a **quantum time** to each running process.

When the process exceeds its quantum time, the OS **interrupts** it and returns it back to the ready queue.

In other words, its status changes from **running** to **ready**.

The process then waits to be selected by the OS to be re-assigned the CPU.

PROCESS STATE (8) – THE STATE DIAGRAM (7) – THE WAITING & READY STATES

A process **waits** for the I/O device in the **device waiting queue**. As soon as it is assigned the device, it performs the I/O operation.

When the process **completes the I/O operation**, for example, gets the value of the input variable from the user, it becomes **ready** to resume execution.

Therefore, the status changes from **wait** to **ready**.

The process then waits in the **ready queue** until it is assigned a CPU by the OS.

PROCESS STATE (9) – NOTES

The state diagram proves that a CPU does not dedicate itself to a processor from the start of its execution till the end.

In fact, an efficient OS aims to make full use of the available resources.

The CPU scheduler guarantees that processes complete execution in an optimum time. This will be studied in details later in this course.

As the number of the processes in the ready queue increases, as there is a less probability that the CPU stays idle.

The state diagram also explains how a CPU executes a single instruction at a time, however, instructions may belong to multiple processes.

In a multiple CPU environment, two or more processes may be dispatched from the ready queue in the same cycle to different processors.

PROCESS CONTROL BLOCK

Each process is represented to the OS by its **Process Control Block (PCB)**.

A PCB contains the following information about the corresponding process:

The process state:

New/Ready/Running/Waiting/Terminated.

The process identifier: The OS assigns a unique number to each process as soon as it is created.

The program counter: This holds the address of the next instruction to be executed.

The CPU registers: The final values produced by the process at an instant of time and stored in the CPU registers. When a process is interrupted by an ISR for example, these values are stored in the PCB. When the process resumes execution, they are restored. We say that the **process state** is saved/restored.

Memory Management Information: The values of the base and limit registers, and the page tables. This are to be explained later in this course.

I/O Status Information: This includes the list of I/O devices allocated to the process & list of open files.

CPU Scheduling Information: This includes the process priority & pointers to scheduling queues.

Accounting Information used for billing. It includes the amount used of CPU, A/C numbers, etc...

process state

process number

program counter

registers

memory limits

list of open files

...