


CHAPTER 3

PROCESS CREATION

gettyimages®

Daniel Grill



PROCESS OPERATIONS

OSs should be able to perform principal operations on processes such as:

- Create a process
- Destroy (kill) a process
- Suspend a process
- Resume a process
- Change the priority of a process
- Block a process
- Wake up a process
- Dispatch a process
- Enable a process to communicate with other processes: this is known as **Inter-Process Communication (IPC)**

PROCESS CREATION (I) – INTRODUCTION

A process may **spawn** or **fork** (create) another new process.


The creating process is called the **parent**. The created one is the **child**.

A child process may itself spawn another process.

This forms a hierarchical structure: a **tree**.

A process may create many children

A process has at most one parent



PROCESS CREATION (2) – PARENT-CHILD RELATIONSHIP

The OS defines the relationship between a parent and its spawned child in three main aspects:

→ **Memory Address Space** – there are two possible scenarios:

1. Both parent and child have the same memory address space
2. Parent and child have different memory address spaces

→ **Mode of Execution** – two scenarios are possible:

1. The process continues to execute concurrently with its child
2. The parent waits until some or all of its children terminate

→ **Resources** – when a parent process spawns a child, the latter needs resources such as CPU time, memory, files, I/O devices, etc... to accomplish its task.

The OS assigns resources to a child process in one of the following scenarios:

1. The child gets its resources directly from the OS
2. A child may be restricted to a subset of the parent's resources.

The second scenario prevents any process from **overloading** the OS by creating too many child processes

PROCESS CREATION (3) – EXAMPLES OF OSs

We'll study the process creation concept in two common operating systems:

→ Unix

→ Windows

PROCESS CREATION (4) – UNIX(I) – INTRODUCTION

In Unix-based OSs, the first created process is called *init*.

init is created as soon as the kernel loads.

Many processes are then spawned directly from *init*.

Each created process has its own identifier *pid*. This is stored in the corresponding PCB. The *init* process is given a *pid* = 1.

The following figure gives an example for a map of processes spawning in Unix:

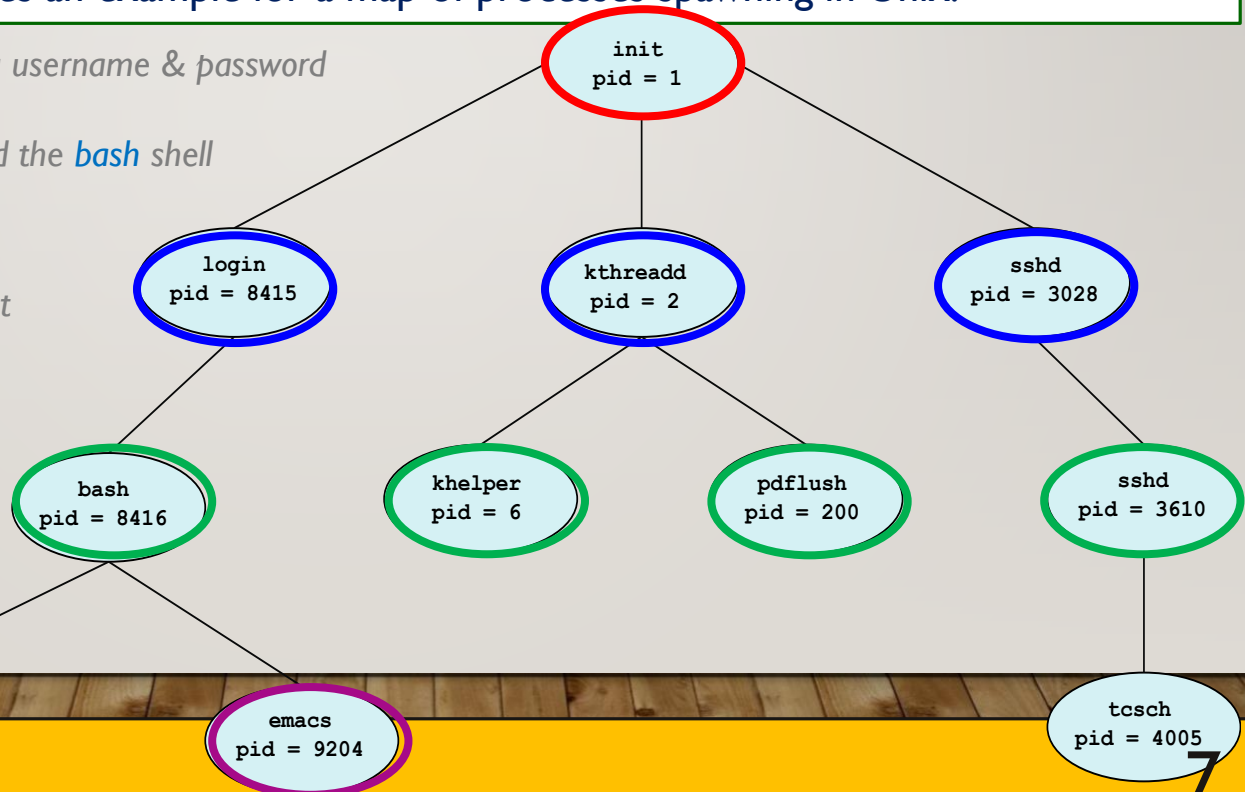
login authenticates the user with a username & password

After authentication, *login* spawned the *bash* shell

Using *bash*, the user created the process *ps* and the *emacs* editor.
ps lists complete information about active processes.

kthreadd is responsible for creating additional processes that perform tasks on behalf of the kernel, such as *khelper* and *pdflush*.

sshd manages clients using *ssh*



PROCESS CREATION (5) – UNIX(2) – MEMORY ADDRESS SPACE (1)

In Unix, a new process is created using the `fork()` system call.

When spawned, the child process “copies” the memory address space of the parent.

`fork()` returns a negative number if there is an error ; a zero for the child process.

The child process overlays its address space using the `exec/p()` command, a version of the `exec()` system call.

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```


PROCESS CREATION (6) – UNIX(3) – MEMORY ADDRESS SPACE (2)

In fact, in UNIX when a child process is created, the OS makes an exact copy of the parent's address space and give it to the child. Therefore, any modifications to the variables made after the *fork()* statement are not reflected in the address space of each other.

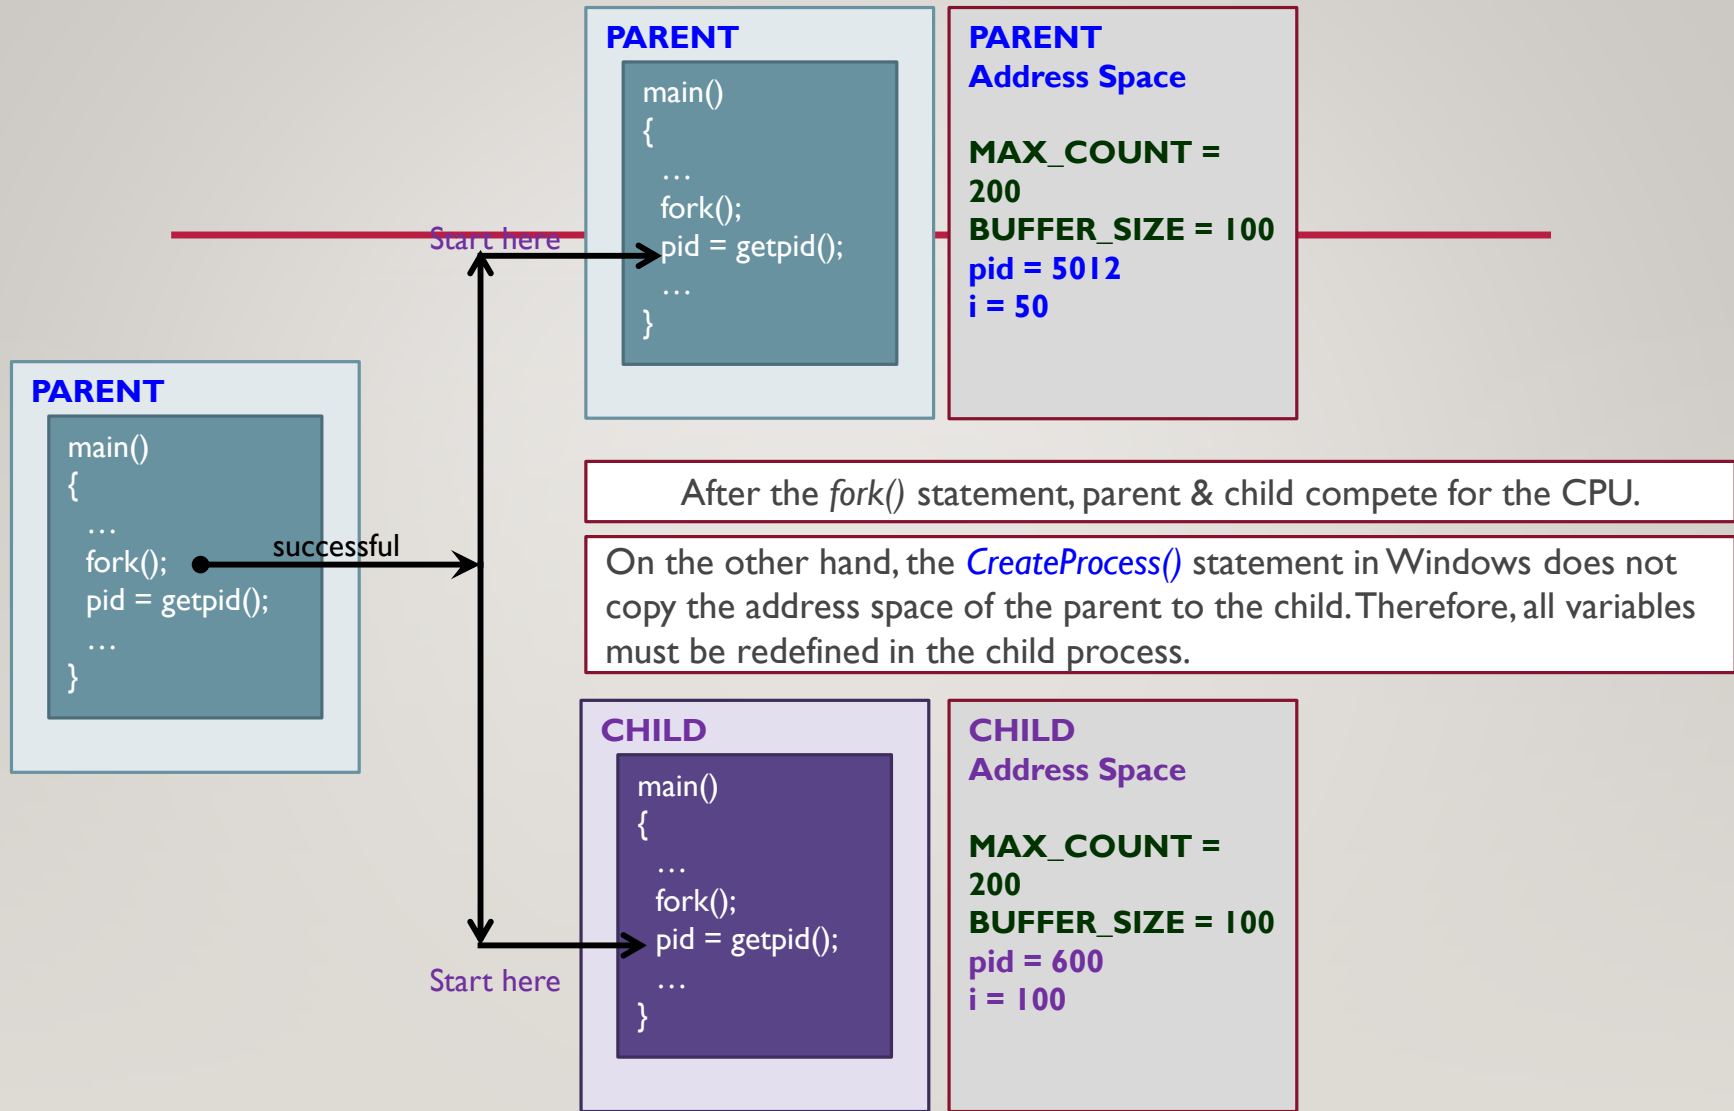
Consider the following example:

```
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#define MAX_COUNT 200;
#define BUFFER_SIZE 100;
void main(void)
{
    pid_t pid;
    int i;
    char buffer[BUFFER_SIZE];
    fork();
    pid = getpid(); //returns the ID of the running process
    for (i=1; i < MAX_COUNT; i++) {
        sprintf (buf, "this line is from pid = %d, value = %d\n", pid, i);
        write (1, buf, strlen(buf)); } //end for
} //end main
```

Values initialized before the *fork()* have the same values in both address spaces (parent & child)

Values modified after the *fork()* are not reflected in the copy's address space

PROCESS CREATION (7) – UNIX(4) – MEMORY ADDRESS SPACE (3)



PROCESS CREATION (6) – UNIX(3) – RESOURCES

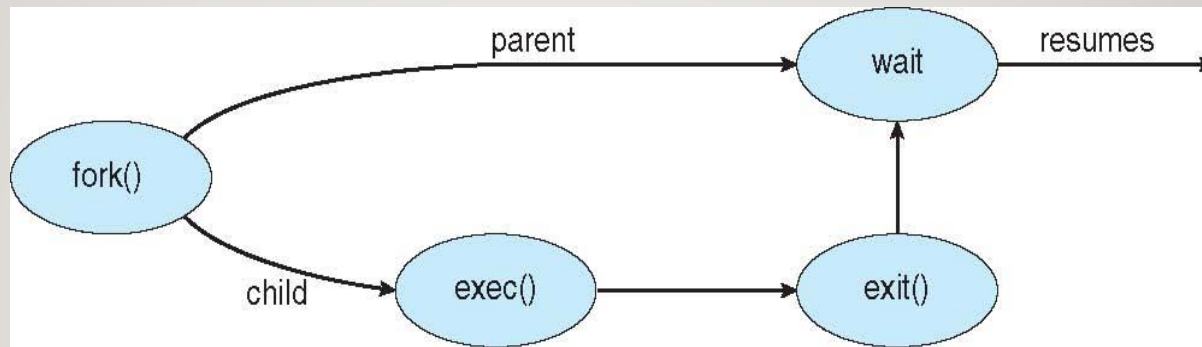
The child process then inherits privileges and scheduling attributes from the parent, as well as certain resources, such as open files.

PROCESS CREATION (7) – UNIX(4) – MODE OF EXECUTION

In Unix, the parent waits for the child process to complete.

This is achieved using the `wait()` command. Refer to the code in slide 8.

When the child completes, then the parent resumes its execution.



```

pid = fork();

if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    return 1;
}
else if (pid == 0) { /* child process */
    execlp("/bin/ls", "ls", NULL);
}
else { /* parent process */
    /* parent will wait for the child to complete */
    wait(NULL);
    printf("Child Complete");
}

return 0;
}
  
```

PROCESS CREATION (8) – WINDOWS(1) – MEMORY ADDRESS SPACE

Windows uses the `CreateProcess()` to create a child process.

In Windows, a child process has a different address space than its parent.

`CreateProcess()` has many parameters.

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

PROCESS CREATION (9) – WINDOWS(2) – RESOURCES

In Windows, the OS assigns resources to child processes directly regardless of those of the parent process.

PROCESS CREATION (10) – WINDOWS(3) – MODE OF EXECUTION

Like Unix, the parent process waits for the completion of its child before resuming execution.

Re-consider the code previously given in slide 11.

```
{
STARTUPINFO si;
PROCESS_INFORMATION pi;

/* allocate memory */
ZeroMemory(&si, sizeof(si));
si.cb = sizeof(si);
ZeroMemory(&pi, sizeof(pi));

/* create child process */
if (!CreateProcess(NULL, /* use command line */
"C:\\WINDOWS\\system32\\mspaint.exe", /* command */
NULL, /* don't inherit process handle */
NULL, /* don't inherit thread handle */
FALSE, /* disable handle inheritance */
0, /* no creation flags */
NULL, /* use parent's environment block */
NULL, /* use parent's existing directory */
&si,
&pi))
{
    fprintf(stderr, "Create Process Failed");
    return -1;
}

/* parent will wait for the child to complete */
WaitForSingleObject(pi.hProcess, INFINITE);
printf("Child Complete");

/* close handles */
CloseHandle(pi.hProcess);
CloseHandle(pi.hThread);
}
```