

PART 4

PROCESS SYNCHRONIZATION - SOFTWARE

gettyimages®

Daniel Grill

→ Producer-consumer Problem

→ Race Condition

→ Critical Section Problem

→ Requirements to the Solution of the CS Problem

→ Peterson's Solution

INTRODUCTION

Simultaneous access to shared data may result in data inconsistency.

Therefore, the execution of cooperating processes that share a logical address space must be organized (**synchronized**) in a way that ensures **data integrity** (or data consistency).

This problem is known as **race condition**:

A situation where several processes access and manipulate the same data simultaneously and the outcome of the execution depends on the particular order in which the access takes place.

PRODUCER-CONSUMER PROBLEM (I) – REVISITED

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Buffer is full → producer waits

An item is produced

Update the position of the *in* pointer

```
item next_consumed;
while (true) {
    while (counter == 0)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next consumed */
}
```

Buffer is empty → consumer waits

An item is consumed

Update the position of the *out* pointer

PRODUCER-CONSUMER PROBLEM (2)

Remember that both code segments of the previous slide are running simultaneously.

They are also sharing a variable; namely, **counter**.

The updated version of the previous slide is as follows:

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

```
while (true) {  
    while (counter == 0)  
        /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

RACE CONDITION – SCENARIO (I)

Although the previous code segments run correctly when run sequentially, they might not run correctly if they are executed simultaneously because the order of execution is unpredictable.

Assume the current value of the variable `counter = 5`.

Consider the following scenario:

P

```
while (true) {
    /* produce an item in next produced */ 1
    while (counter == BUFFER_SIZE) ;
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++; 3
}
```

- At the start, `counter = 5`
- Segment (1) of the producer is executed
- Segment (2), the consumer, is assigned to the CPU before counter is incremented →

- `counter = 4`
- Segment (3) of the producer executes → `counter = 5;`
- In fact, there are 6 items in the buffer

C

```
while (true) {
    while (counter == 0) 2
        /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```

- `counter = 5`
- After the consumer executes, `counter = 4`
- In fact, there are 5 items in the buffer

RACE CONDITION – SCENARIO (2)

Now consider the following sequence:

At the start, the variable `counter = 5`.

P

```
while (true) {
    /* produce an item in next produced */
    while (counter == BUFFER_SIZE);
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

2

- `counter = 5`
- After the producer executes, `counter = 6`
- In fact, there are 5 items in the buffer

C

```
while (true) {
    while (counter == 0)
        /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```

1

3

- Initially, `counter = 5`
- Segment (1) of the consumer executes
- Segment (2), the producer, is assigned to the CPU before counter is decremented →

- `counter = 6`
- Segment (3) of the consumer executes → `counter = 5`
- In fact, there are 4 items in the buffer

RACE CONDITION – SCENARIO (3)

Now, consider the following sequence:

Assume the variable `counter = 5`.

P

```
while (true) {
    /* produce an item in next produced */
    while (counter == BUFFER_SIZE);
        /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

1

- Initially, `counter= 5`
- Segment (1), the whole producer code, is executed → `counter= 6`.
- In fact, there are 6 items in the buffer

C

```
while (true) {
    while (counter == 0)
        /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```

2

- `counter= 6`
- After Segment (2), the consumer, executes → `counter= 5`
- In fact, there are 5 items in the buffer

RACE CONDITION – SCENARIO (4)

Now, consider the following sequence:

Assume the variable `counter = 5`.

P

```
while (true) {
    /* produce an item in next produced */
    while (counter == BUFFER_SIZE) ;
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

2

- `counter= 4`
- Segment 2, the producer code executes → `counter= 5`
- In fact, there are 5 items in the buffer

C

```
while (true) {
    while (counter == 0)
        /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```

1

- Initially, `counter= 5`
- Segment (1), the whole consumer, executes → `counter= 4`
- In fact, there are 4 items in the buffer

RACE CONDITION – LOW-LEVEL STATEMENTS

The high-level language statements `counter++` & `counter--` are implemented in machine language on three steps each as follows:

`counter++`

```
R1 = counter
R1 = R1 + 1
counter = R1
```

`counter--`

```
R2 = counter
R2 = R2 - 1
counter = R2
```

Where R1 and R2 are the local register on the CPU.

The simultaneous execution of both statements is equivalent to a sequential execution in which the low-level statements are interleaved in some arbitrary order.

However, the order of high-level statements is preserved.

Example of the interleaving sequential low-level statements for the execution of `counter++` and `counter--` simultaneously is shown below. Assume that the initial value of counter is 5.

S0: producer executes	<code>R1 = counter</code>	{R1 = 5}
S1: producer executes	<code>R1 = R1 + 1</code>	{R1 = 6}
S2: consumer executes	<code>R2 = counter</code>	{R2 = 5}
S3: consumer executes	<code>R2 = R2 - 1</code>	{R2 = 4}
S4: producer executes	<code>counter = R1</code>	{counter = 6}
S5: consumer executes	<code>counter = R2</code>	{counter = 4}

In fact, there are five items in the buffer (not 4) by the end of execution of S5.

In another run, we might have S5 executing before S4. Thus the incorrect value would be 6 instead of 5.

CRITICAL SECTION PROBLEM (I) - OVERVIEW

For two cooperative processes P_i and P_j , sharing a common variable *shared*, the code may be divided into two parts:

- The part of the code of P_i (or P_j), in which the common variable *shared* is used. This is called the critical section (C).
- The rest of the code of P_i and P_j that does not use any common variable. This is called the remainder section (R).

In order to avoid the race condition problem, P_i and P_j should be synchronized.

Synchronization is performed by prohibiting the simultaneous execution of the critical sections C_i and C_j , corresponding to the processes P_i and P_j .

Therefore, for a cooperative process P_i , before entering its own critical section C_i , it should ensure that no other cooperative process is using the common variable *shared*.

In addition, a process P_i should announce its exit from C_i ; ie. P_i does not use any more the common variable *shared* so that it can be used by the other process P_j .

In other words, two announcements are necessary for each cooperative process P_i when using a common variable:

- 1) P_i makes an announcement before entering its critical section C_i .
This prohibits any other cooperative process P_j to enter C_j .
- 2) P_i makes an announcement as soon as it exits from its critical section C_i .
This allows any other cooperative process P_j to enter C_j .

In other words, P_i captures the common variable *shared* in the first announcement, and releases it in the second announcement.

CRITICAL SECTION PROBLEM (2) – PROTOCOL

For a set of cooperative processes, the critical section is defined to be the segment of code in which the process may be changing shared variables, updating a table, writing a file, etc...

The rule to avoid the race condition is that when one process P_i is executing in its critical section C_i , no other cooperative process P_j is allowed to execute in its own critical section C_j .

In other words, no two cooperative processes are allowed to execute in their critical sections at the same time.

The critical section problem aims to design a protocol for the cooperative processes so that they fulfill the above rule.

The protocol of the critical section problem consists of four main parts:

→ 1) The entry section

This is the code segment in which a cooperative process P_i requests permission to enter its critical section C_i .

→ 2) The critical section

This is the code segment in which a process P_i uses shared variables.

→ 3) The exit section

This is the code segment in which a cooperative process P_i releases the shared variable(s). This is done when P_i exits C_i .

→ 4) The remainder section

This is the code segment in which a process P_i does not use any shared variables.

CRITICAL SECTION PROBLEM (3) – PSEUDO-CODE

 P_i

do {

entry section

critical section

exit section


remainder section

} while (true);

REQUIREMENTS TO THE SOLUTION OF THE CS PROBLEM

Many solutions have been proposed to the critical section (CS) problem.

An efficient solution to the CS problem should satisfy the following three requirements:



REQUIREMENTS TO THE SOLUTION OF THE CS PROBLEM – MUTUAL EXCLUSION

An efficient solution should ensure **mutual exclusion** between cooperative processes.

If a process P_i is executing in its critical section C_i , then no other processes P_j can be executing in their corresponding critical section C_j at the same time.

In other words, one critical section is executing at a time.

Note that we may have the following scenario:

- A process P_1 shares a variable $sharedI-2$ with a process P_2
- The process P_1 shares another variable $sharedI-3$ with a process P_3
- In such case there are 4 critical sections (P_1, P_2) , (P_1, P_3) , (P_2, P_1) , and (P_3, P_1)
- The solution of the mutual exclusion problem does not allow (P_1, P_2) and (P_2, P_1) to execute simultaneously.
- Also, (P_1, P_3) and (P_3, P_1) are not allowed to execute simultaneously.
- However, (P_1, P_2) and (P_3, P_1) for example may execute simultaneously since they are working on different shared variables; namely, $sharedI-2$ and $sharedI-3$ respectively.

REQUIREMENTS TO THE SOLUTION OF THE CS PROBLEM – PROGRESS

An efficient solution should also ensure the **progress** of all cooperative processes.

If two (or more) cooperative processes **P1** and **P2** require permission to enter their respective critical sections sharing the same variable **shared** at a time **T1**, then two conditions should be satisfied to ensure the progress requirement:

Only **P1** and **P2**, are given the right to decide which process (**P1** or **P2**) is given the permission first to enter its critical section. Other processes sharing the same variable but running in their Remainder Section are not involved in this decision.

The negotiation decision should be made as soon as possible.

Consider the following example for three cooperative processes **P1**, **P2** and **P3** sharing the same variable:

At some point of time **T0**, there is no process executing in its CS.

At time **T1**, **P1** and **P2** wish to enter their CS.

Only **P1** and **P2** negotiate to decide which of them enters the CS.

Since **P3** is executing in its Remainder Section, it is not involved in the decision.

The decision should be taken as soon as possible.


Assume that the decision is given to the favor of **P1**, then **P2** waits.

REQUIREMENTS TO THE SOLUTION OF THE CS PROBLEM – BOUNDED WAITING

An efficient solution should also ensure **bounded waiting** for any cooperative process

By bounded waiting, it is meant that there is a limit (bound) in the number of times in which a process makes a request to enter its CS before it is given permission to enter its CS.

Consider the following example for five cooperative processes *P1*, *P2*, *P3*, *P4* and *P5* sharing the same variable. Assume that the *bound = 3 times*.

- 
- Initially, *bound = 0*, for all processes.
 - P3* & *P4* wish to enter their CS. They negotiate, and the decision is in favor of *P4*
 - P4* enters its CS, and *bound* is incremented (*= 1*) for *P3*.
 - Later, *P3* & *P1* wish to enter their CS. The decision is in favor of *P1*.
 - P1* enters its CS and the bound of *P3* is incremented (*bound=2*)
 - Later, *P3* and *P2* wish to enter their CS. The decision is in favor of *P2*.
 - P2* enters its CS, and the bound of *P3* is incremented (*bound=3*).
 - Later, *P3* and *P5* wish to enter their CS: since the bound of *P3* reached its maximum allowed limit, then *P3* enters its CS without negotiation.

In other words, no process should wait infinitely to enter its CS.

The figure in the next slide illustrates this example.

BOUNDED WAITING – EXAMPLE

(1) P3 & P4 compete for the CS

Decision in favor of P4, B(P3)++

(2) P3 & P1 compete for the CS

Decision in favor of P1, B(P3)++

(3) P3 & P2 compete for the CS

Decision in favor of P2, B(P3)++



(4) P3 & P5 compete for the CS

P3 reached threshold (maximum bound=3)

P3 is guaranteed to enter its CS

B(P3) is reset to 0

B(P5)++




PETERSON'S SOLUTION

This is a classic software-based solution to the CS problem.

Peterson's solution is restricted to two cooperative processes only.

With two cooperative processes, *P0* and *P1*, Peterson's solution is designed so that they alternate between the execution of their Critical and Remainder Sections.

Peterson's solution requires the two processes to share two data items; namely:



PETERSON'S SOLUTION – THE TURN VARIABLE

The variable *turn* indicates whose turn it is to enter the CS: *P0* or *P1*.

Therefore, *turn* equals either to 0 or 1, since we deal with two processes only.

When *turn* = 0, it is the turn of *P0* to enter its CS.

When *turn* = 1, it is the turn of *P1* to enter its CS.

The next turn may also be expressed as $1 - i$, where *i* is the number of the current process.

It might happen that both processes try to enter their CS at the same time.

- In this case, both processes update the value of *turn* at the same time
- However, only one of these values will last and the other is overwritten
- We cannot predict which process will enter its CS first.

PETERSON'S SOLUTION – THE FLAG ARRAY

The *flag* array indicates if a process is ready to enter its critical section.

If *flag[0] = true*, then *P0* wishes to enter its CS.

If *flag[1] = true*, then *P1* wishes to enter its CS.

If both processes wish to enter their CS, then the variable *turn* decides which process will enter the CS.

PETERSON'S SOLUTION – THE CODE

The following figure depicts the code segments for two cooperative processes, *P0* and *P1*, that are running simultaneously.

Po	P1
<pre> do { flag[0] = true; turn = 1; while (flag[1] && turn == 1); critical section flag[0] = false; remainder section } while (true); </pre>	<pre> do { flag[1] = true; turn = 0; while (flag[0] && turn == 0); critical section flag[1] = false; remainder section } while (true); </pre>

Does Peterson's solution fulfill the requirements of the solution to the CS?

- The while loop ensures **mutual exclusion**.
- **Progress** is satisfied since turns are assigned to *P0* and *P1* alternatively.
- For the same reason, no process waits infinitely → **bounded waiting** is satisfied.

Since the three requirements are fulfilled, then Peterson's solution is considered an efficient to the CS problem..

Although Peterson's solution is designed for two cooperative processes; the algorithm can be updated to fit three or more cooperative processes.