

Ch.5 – Process Synchronization

- Race Condition: several processes access and manipulate the same data concurrently, outcome depends on which order each access takes place.
- Each process has critical section of code, where it is manipulating data
 - To solve critical section problem each process must ask permission to enter critical section in entry section, follow critical section with exit section and then execute the remainder section
 - Especially difficult to solve this problem in preemptive kernels
- Peterson's Solution: solution for two processes
 - Two processes share two variables: int **turn** and Boolean **flag[2]**
 - **turn**: whose turn it is to enter the critical section
 - **flag**: indication of whether or not a process is ready to enter critical section
 - `flag[i] = true` indicates that process P_i is ready
 - Algorithm for process P_i :

```
do {
    flag[i] = TRUE;
    turn = j;
    while (flag[j] && turn == j)
        critical section
    flag[i] = FALSE;
    remainder section
} while (TRUE);
```

- Modern machines provide atomic hardware instructions: Atomic = non-interruptible
- Solution using Locks:

```
do {
    acquire lock
    critical section
    release lock
    remainder section
} while (TRUE);
```

- Solution using Test-And-Set: Shared boolean variable lock, initialized to FALSE

```
boolean TestAndSet (boolean *target){
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

```
do {
    while ( TestAndSet (&lock ))
        ; // do nothing
    // critical section
    lock = FALSE;
    // remainder section
} while (TRUE);
```

- Solution using Swap: Shared bool variable lock initialized to FALSE; Each process has local bool variable key

```
void Swap (boolean *a, boolean *b){
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

```
do {
    key = TRUE;
    while ( key == TRUE)
        Swap (&lock, &key );
    // critical section
    lock = FALSE;
    // remainder section
} while (TRUE);
```

- Semaphore: Synchronization tool that does not require busy waiting (it has a busy-waiting version, by the non-busy-waiting version is the most common)
 - Standard operations: `wait()` and `signal()` ← these are the only operations that can access semaphore S
 - Can have counting (unrestricted range) and binary (0 or 1) semaphores
- Deadlock: Two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes (most OSes do not prevent or deal with deadlocks)
 - Can cause starvation and priority inversion (lower priority process holds lock needed by higher-priority process)