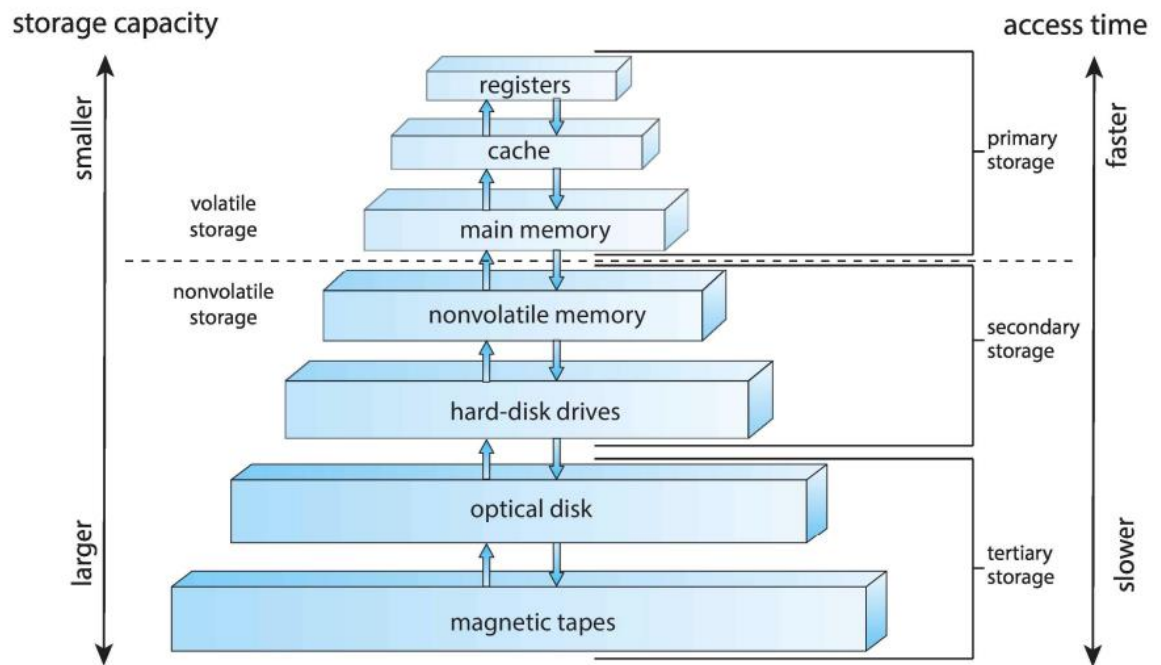


Chapter 1

- 1- **Operating system** is a resource allocator and control program making efficient use of HW and managing execution of user programs
- 2- Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- 3- The one program running at all times on the computer is the **kernel**, part of the operating system
- 4- a **system program** (ships with the operating system, but not part of the kernel)
- 5- an **application program**, all programs not associated with the operating system
- 6- **middleware** – a set of software frameworks that provide additional services to application developers such as databases, multimedia, graphics
- 7- One or more CPUs, device controllers connect through common **bus** providing access to shared memory
- 8- Each device controller type has an operating system **device driver** to manage it
- 9- Device controller informs CPU that it has finished its operation by causing an **interrupt**
- 10- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- 11- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- 12- An operating system is **interrupt driven**
- 13- **bootstrap program** is loaded at power-up or reboot
- 14- **bootstrap program** Typically stored in ROM or EPROM, generally known as firmware
- 15- **System call** – request to the OS to allow user to wait for I/O completion
- 16- **Device-status table** contains entry for each I/O device indicating its type, address, and state
- 17- **Main memory** – only large storage media that the CPU can access directly
- 18- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
- 19- **Hard Disk Drives (HDD)** – rigid metal or glass platters covered with magnetic recording material
- 20- Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
- 21- The **disk controller** determines the logical interaction between the device and the computer
- 22- **Non-volatile memory (NVM)** devices– faster than hard disks, nonvolatile
- 23- Storage systems organized in hierarchy **Speed, Cost then Volatility**
- 24- **Caching** – copying information into faster storage system; main memory can be viewed as a cache for secondary storage
- 25- **Device Driver** for each device controller to manage I/O, it provides uniform interface between controller and kernel



- 26-
- 27- Multiprocessors systems growing in use and importance, Also known as parallel systems, tightly-coupled systems
- 28- Multiprocessors Advantages include: 1. Increased throughput 2. Economy of scale 3. Increased reliability – graceful degradation or fault tolerance
- 29- Asymmetric Multiprocessing – each processor is assigned a specific task
- 30- Symmetric Multiprocessing – each processor performs all tasks
- 31- Usually sharing storage via a storage-area network (SAN)
- 32- Asymmetric clustering has one machine in hot-standby mode
- 33- Symmetric clustering has multiple nodes running applications, monitoring each other
- 34- Some clusters are for high-performance computing (HPC), Applications must be written to use parallelization
- 35- Some have distributed lock manager (DLM) to avoid conflicting operations
- 36- Multiprogramming (Batch system) needed for efficiency
- 37- Timesharing (multitasking) is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating interactive computing.
- 38- Response time should be < 1 second
- 39- Each user has at least one program executing in memory -> process
- 40- If several jobs ready to run at the same time -> CPU scheduling
- 41- If processes don't fit in memory, swapping moves them in and out to run
- 42- Virtual memory allows execution of processes not completely in memory
- 43- Dual-mode operation allows OS to protect itself and other system components
- 44- Single-threaded process has one program counter specifying location of next instruction to execute
- 45- Multiprocessor environment must provide cache coherency in hardware such that all CPUs have the most recent value in their cache

- 46- **Protection** – any mechanism for controlling access of processes or users to resources defined by the OS
- 47- **Security** – defense of the system against internal and external attacks
- 48- **Emulation** used when source CPU type different from target type (i.e. PowerPC to Intel x86)
- 49- **Virtualization** – OS natively compiled for CPU, running guest OSes also natively compiled
- 50- **Network** is a communications path, **TCP/IP** most common– **Local Area Network (LAN)** – **Wide Area Network (WAN)** – **Metropolitan Area Network (MAN)** – **Personal Area Network (PAN)**
- 51- **Network Operating System** provides features between systems across network
- 52- **Hash function** can create a **hash map**
- 53- **Bitmap** – string of n binary digits representing the status of n items
- 54- **Portals** provide web access to internal systems
- 55- **Network computers (thin clients)** are like Web terminals
- 56- Mobile computers interconnect via **wireless networks**
- 57- Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks
- 58- Leaders are **Apple iOS** and **Google Android**
- 59- **Compute-server system** provides an interface to client to request services (i.e., database)
- 60- **File-server system** provides interface for clients to store and retrieve files
- 61- **Public cloud** – available via Internet to anyone willing to pay
- 62- **Private cloud** – run by a company for the company's own use
- 63- **Hybrid cloud** – includes both public and private cloud components
- 64- Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
- 65- Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
- 66- Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)

Chapter 4

- 1- **Process creation** is heavy-weight (time consuming and resource intensive) while thread creation is light-weight (because threads share the code, data and OS resources)
- 2- If the **web-server** process is **multithreaded**, the server will create a separate thread that listens for client requests
- 3- When a request is made, rather than creating another process, the server creates a new thread to service the request and resume listening for additional requests.
- 4- If the web server ran as a traditional single-threaded process, it would be able to service only **one client at a time, and a client might have to wait a very long time** for its request to be serviced.
- 5- **Kernels** are generally **multithreaded**
- 6- Processes can only share resources through techniques such as **shared memory** and **message passing**.
- 7- A **single-threaded process** can run on only one processor, regardless how many are available.
- 8- **Multicore or multiprocessor** systems putting pressure on programmers, (**to make better use of the multiple computing cores**)
- 9- **Parallelism** implies a system can perform more than one task simultaneously
- 10- **Concurrency** allows more than one task to make progress
- 11- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
- 12- **Task parallelism** – distributing threads across cores, each thread performing unique operation
- 13- Three primary thread libraries: **POSIX Pthreads**, **Windows threads** and **Java threads**.
- 14- **Thread library** provides programmer with API for creating and managing threads
- 15- Invoking a function in the API for the library typically results in a system call to the kernel.
- 16- **Asynchronous threading**, once the parent creates a child thread, the parent resumes its execution

- 17- **Synchronous threading** occurs when the parent thread creates one or more children and then must wait for all of its children to terminate before it resumes
- 18- **POSIX stands for Portable Operating System Interface**, is a family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems.
- 19- Any data declared globally—that is, declared outside of any function—are shared among all threads belonging to the same process.
- 20- Because Java has no notion of global data, access to shared data must be explicitly arranged between threads.
- 21- There are **two techniques** for creating threads in a Java program. One approach is **to create a new class that is derived from the Thread class and to override its run() method**. An alternative—and **more commonly used**—**technique is to define a class that implements the Runnable interface**.
- 22- If two or more threads are to share data in a Java program, the sharing occurs by passing references to the shared object to the appropriate threads.
- 23- **parallel regions** – blocks of code that can run in parallel
- 24- if a thread invokes **the exec()** system call, the program specified in the parameter to exec() will replace the entire process—including all threads.
- 25- **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- 26- A **signal handler** is used to process signals.
- 27- Every signal has **default handler** that kernel runs when handling signal
- 28- **User-defined signal** handler can override default
- 29- A signal may be received either **synchronously** or **asynchronously**
- 30- Typically, an **asynchronous** signal is sent to another process.
- 31- **synchronous** signal include: **illegal memory access** and **division by 0**
- 32- Thread to be canceled is **target thread**.
- 33- Cancellation only occurs when thread reaches **cancellation point**

- 34- **Thread-local storage (TLS)** allows each thread to have its own copy of data
- 35- **Lightweight process (LWP)** Typically use an intermediate data structure between user and kernel threads
- 36- **upcalls** - a communication mechanism from the kernel to the upcall handler in the thread library
- 37- The register set, stacks, and private storage area are known as the **context of the thread**

Chapter 6

- 1- Each process has **critical section** segment of code
- 2- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**
- 3- **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
- 4- **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
- 5- **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
- 6- **Preemptive** – allows preemption of process when running in kernel mode
- 7- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
- 8- **Mutex Locks** requires **busy waiting**. This lock called spinlock
- 9- **Counting semaphore** – integer value can range over an unrestricted domain
- 10- **Binary semaphore** and **mutex lock** – integer value can range only between 0 and 1
- 11- **block** – place the process invoking the operation on the appropriate waiting queue
- 12- **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- 13- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Chapter 9

- 1- Main memory can take many cycles, causing a **stall**
- 2- **Cache** sits between main memory and CPU registers
- 3- We can provide protection by using a pair of **base** and **limit registers** define the logical address space of a process
- 4- **Compile time**: If memory location known a priori, absolute code can be generated; must **recompile** code if starting location changes
- 5- **Load time**: Must generate **relocatable code** if memory location is not known at compile time
- 6- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
- 7- **Logical address** – generated by the CPU; also referred to as virtual address
- 8- **Physical address** – address seen by the memory unit
- 9- **Logical address space** is the set of all logical addresses generated by a program
- 10- **Physical address space** is the set of all physical addresses generated by a program
- 11- Dynamically linked libraries are system libraries (such as language subroutine libraries) that are linked to user programs when the programs are run.
- 12- **Static linking** – system libraries and program code combined by the loader into the binary program image
- 13- Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- 14- System also known as **shared libraries**
- 15- A library may be **replaced by a new version**, and all programs that reference the library will automatically use the new version.
- 16- Main memory usually into two **partitions**, **OS** and **user**
- 17- **Hole** – block of available memory; holes of various size are scattered throughout memory
- 18- **First-fit**: Allocate the first hole that is big enough
- 19- **Best-fit**: Allocate the smallest hole that is big enough; must search entire list, unless ordered by size
- 20- **Worst-fit**: Allocate the largest hole; must also search entire list
- 21- First-fit and best-fit better than worst-fit in terms of **speed** and **storage utilization**

- 22- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- 23- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- 24- Reduce external fragmentation by **compaction**
- 25- Divide physical memory into fixed-sized blocks called **frames**
- 26- Divide logical memory into blocks of same size called **pages**
- 27- Set up a **page table** to translate logical to physical addresses
- 28- **Page number (p)** – used as an index into a page table which contains base address of each page in physical memory
- 29- **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit
- 30- **Page-table base register (PTBR)** points to the page table
- 31- **Page-table length register (PTLR)** indicates size of the page table
- 32- The two memory access problem can be solved by the use of a special fast-lookup hardware cache **called translation look-aside buffers (TLBs)** (also called **associative memory**)
- 33- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
- 34- Some entries can be **wired down** for permanent fast access
- 35- Valid-invalid bit attached to each entry in the page table: “**valid**” indicates that the associated page is in the process’ logical address space, and is thus a legal page “**invalid**” indicates that the page is not in the process’ logical address space
- 36- Variation for 64-bit addresses is **clustered page tables**
- 37- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
- 38- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- 39- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- 40- System maintains **a ready queue** of ready-to-run processes which have memory images on disk

41- Android terminates apps if low free memory, but first writes **application state** to flash for fast restart

42- 32-bit address limits led Intel to create **page address extension (PAE)**, allowing 32-bit apps access to more than 4GB of memory space, Top two bits refer to a **page directory pointer table**

43-